

```
In [1]: 1 matplotlib inline
2 from __future__ import print_function
3 import argparse
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
7 import torch.optim as optim
8 import torchvision
9 from torch.utils.data import Dataset
10 from torchvision import datasets, transforms
11 from torch.autograd import Variable
12
13 import matplotlib as mpl
14 import matplotlib.pyplot as plt
15 import seaborn as sns
16 import numpy as np
17 import seaborn as sns
18 import pandas as pd
19 from tqdm import tqdm, tqdm_notebook, tqdm
20 import math
21 import random
22 from PIL import Image
23 import os
24 from collections import OrderedDict
```

Preprocess

```
In [2]: 1 DataDir = './'
2 SAPARATELINE = '.' * 100
3 SAPARATE = lambda: print(SAPARATELINE)
4 data = pd.read_csv('train.csv')
5
```

The shape of 'tr'

```

0         Unnamed: 0  id          left          right \
1         0         Data/01.90+00000_left.jpg  Data/01.90+00000_right.jpg
2         1         Data/01.90+00001_left.jpg  Data/01.90+00001_right.jpg
3         2         Data/01.90+00002_left.jpg  Data/01.90+00002_right.jpg
4         3         Data/01.90+00003_left.jpg  Data/01.90+00003_right.jpg
5         4         Data/01.90+00006_left.jpg  Data/01.90+00006_right.jpg

        top_label          left_box          right_box          top_box
0  Data/01.90+00000_top.jpg  90+ ['142', '37', '563', '431']  90+ ['142', '37', '563', '431']
1  Data/01.90+00001_top.jpg  90+ ['93', '35', '555', '442']  90+ ['93', '35', '555', '442']
2  Data/01.90+00002_top.jpg  90+ ['100', '59', '548', '437']  90+ ['100', '59', '548', '437']
3  Data/01.90+00003_top.jpg  90+ ['89', '67', '520', '433']  90+ ['89', '67', '520', '433']
4  Data/01.90+00006_top.jpg  90+ ['116', '77', '520', '426']  90+ ['116', '77', '520', '426']

        right_box          top_box
0  ['104', '22', '567', '442']  ['125', '76', '499', '408']
1  ['120', '45', '522', '420']  ['174', '101', '533', '427']
2  ['119', '70', '546', '426']  ['167', '94', '498', '441']
3  ['142', '66', '535', '428']  ['173', '108', '526', '456']
4  ['118', '19', '589', '445']  ['105', '95', '496', '444']

In [3]: 1  random_index = np.random.randint(0,data.index.size)
        2  # print('Label:\t',data['label'][random_index])
        3
        4  random_apple_info = data.loc[random_index]
        5  img_left_pos = random_apple_info['left_box'][(2:-2).split(",")
        6  img_top_pos = random_apple_info['top_box'][(2:-2).split(",")
        7  img_right_pos = random_apple_info['right_box'][(2:-2).split(",")
        8
        9  print('Label:\t',data['label'][random_index])
        10 print('Left:\t{:30}TopPos: ({:3},{:3}),({:3},{:3})'
        11       .format(random_apple_info.at['left'],img_left_pos[0],img_left_pos[1],img_left_pos[2],img_left_pos[3]))
        12 print('Top:\t{:30}TopPos: ({:3},{:3}),({:3},{:3})'
        13       .format(random_apple_info.at['top'],img_top_pos[0],img_top_pos[1],img_top_pos[2],img_top_pos[3]))
        14 print('Right:\t{:30}TopPos: ({:3},{:3}),({:3},{:3})'
        15       .format(random_apple_info.at['right'],img_right_pos[0],img_right_pos[1],img_right_pos[2],img_right_pos[3]))
        16

Label:      80+
Left:  Data/02.80+00583_left.jpg          Pos: (118,60 ),(456,366)
Top:   Data/02.80+00583_top.jpg           Pos: (215,153),(490,419)
Right: Data/02.80+00583_right.jpg         Pos: (169,138),(450,381)

In [4]: 1  pii_img_left = Image.open(os.path.join(DataDir, random_apple_info.at['left'])).crop((int(img_left_pos[0]),
        2
        3
        4
        5
        6
        7
        8
        9
        10
        11
        12
        13
        14
        15
        16
        17
        18
        19
        20
        21
        22
        23
        24
        25
        26
        27
        28
        29
        30
        31
        32
        33
        34
        35
        36
        37
        38
        39
        40
        41
        42
        43
        44
        45
        46
        47
        48
        49
        50
        51
        52
        53
        54
        55
        56
        57
        58
        59
        60
        61
        62
        63
        64
        65
        66
        67
        68
        69
        70
        71
        72
        73
        74
        75
        76
        77
        78
        79
        80
        81
        82
        83
        84
        85
        86
        87
        88
        89
        90
        91
        92
        93
        94
        95
        96
        97
        98
        99

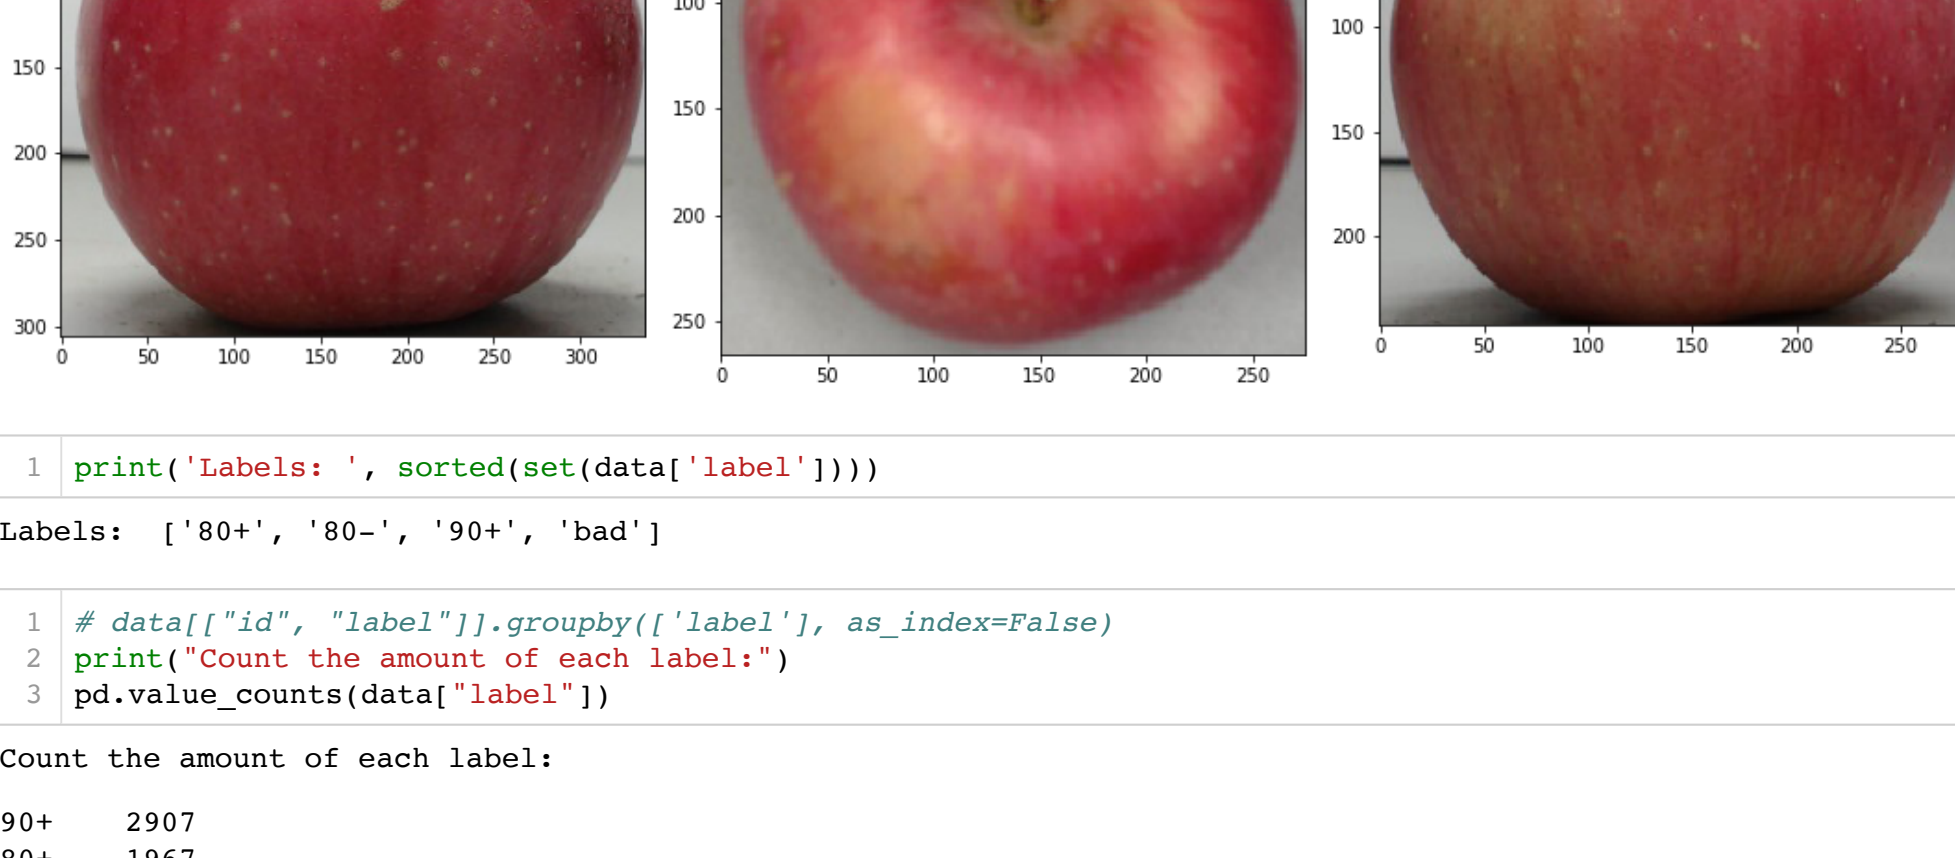
```

```

6 pil_img_right = Image.open(os.path.join(DataDir, random_apple_info[0], "right.jpg"))
7
8 int(img_right_pos[1]),
9 int(img_right_pos[3])
10 pil_img_top = Image.open(os.path.join(DataDir, random_apple_info.at[1] "top.jpg"))
11 crop((int(img_top_pos[0]),
12 int(img_top_pos[1]),
13 int(img_top_pos[2]),
14 int(img_top_pos[3]))
15
16 # pil_im = Image.open('1.jpg').convert('L') #灰度操作
17
18 plt.figure(figsize=(5 * 3, 5))
19
20 ax1 = plt.subplot('131')
21 plt.imshow(pil_img_left)
22 ax1.set_title("Left")
23
24 ax2 = plt.subplot(132)
25 plt.imshow(pil_img_top)
26 ax2.set_title("top")
27
28 ax3 = plt.subplot(133)
29 plt.imshow(pil_img_right)
30 ax3.set_title("right")
31
32 plt.tight_layout()
33 plt.show()

```

The figure displays three subplots arranged horizontally, each showing a different view of an apple. The subplots are labeled 'Left', 'top', and 'right' above them. Each subplot has a vertical axis on the left with tick marks at 0 and 50. The 'Left' subplot shows the left side of the apple, the 'top' subplot shows the top of the apple, and the 'right' subplot shows the right side of the apple. The apple is red with some yellow and green areas.



```
80-      441
bad      396
Name: label, dtype: int64
```

Get mean and std of every channel

This is only a example, just show how to get mean and std from one picture.

Notice: In training phase, your mean and std must calculate from the whole training set.

```
In [7]: 1 pic = np.array(pil_img_left)
2 print("pic.shape:",pic.shape)
3
4 pic_mean = np.mean(pic/255, axis = (0,1),keepdims = True)
5 print("pic_mean:",pic_mean)
```

```
10 pic_std = np.std(pic,axis = (0,1),keepd
```

```
11 print( pic_std: ,pic_std)
12
```

```
14 print("pic_std:",pic_std)
15
```

```
pic.shape: (306, 338, 3)
pic_mean: [[[[0.59334833 0.34683197 0.36197992]]]
pic_mean: [[0.59334833 0.34683197 0.36197992]]
pic_std: [[[[47.27379549 68.99296688 66.41355196]]]
pic_std: [[47.27379549 68.99296688 66.41355196]]
```

start

```
In [8]: 1 class AppleSet(Dataset):
```

4
5
6

```

# 7         'classes' = ['90+', '80+', '80-', 'bad']
# 8         if self.train:
# 9             self.data = pd.read_csv(os.path.join(self.root, 'train.csv'))
# 10            print(pd.value_counts(self.data['label']))
# 11         else:
# 12             self.data = pd.read_csv(os.path.join(self.root, 'test.csv'))
# 13            print(pd.value_counts(self.data['label']))
# 14
# 15         def __len__(self):
# 16             return self.data.index.size
# 17
# 18         def __getitem__(self, index):
# 19             self.classes = ['90+', '80+', '80-', 'bad']
# 20
# 21             data_item = self.data.iloc[index]
# 22             apple_item_path = data_item['left']
# 23             apple_item_img = Image.open(os.path.join(self.root, apple_item_path))
# 24             apple_img_box_list = data_item['left_box'][:2].split(", ")
# 25             apple_item_img = apple_item_img.crop((int(apple_img_box_list[0]),int(apple_img_box_list[1]),int(
# 26
# 27             if self.transform:
# 28                 apple_item_img = self.transform(apple_item_img)
# 29             target = data_item['label']
# 30
# 31             return apple_item_img, self.classes.index(target), data_item['id'], data_item['label']
# 32

```

```

In [9]: 1 N_CLASSES = 4
2 IMAGE_SIZE = 227
3 BATCH_SIZE = 16
4
5 use_cuda = torch.cuda.is_available()
6 if (use_cuda):
7     print("Great, you have a GPU!")
8 else:
9     print("Life is short -- consider a GPU!")
10 device = torch.device("cuda:2" if use_cuda else "cpu")
11 print('Your device will be: ', device)

```

Great, you have a GPU!

Your device will be: cuda:2

```
In [10]: 1 transform = transforms.Compose([
2         # transforms.CenterCrop(480), #???
3         transforms.Resize((IMAGE_SIZE, IMAGE_SIZE)),
4         transforms.RandomHorizontalFlip(), #???
5         transforms.ToTensor(),
6         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
7     ])
8
9 train_data = AppleSet(root = DataDir, train = True, transform = transform)
10 test_data = AppleSet(root = DataDir, train = False, transform = transform)
```

```
12 print(len(train_data))
13 print(len(test_data))

90+ 2907
80+ 1967
80- 441
bad 396
Name: label, dtype: int64
90+ 738
80+ 497
80- 118
bad 75
```

```

      148: label, dtype: int64
      149: 5711
      150: 1428

In [11]: 1 print('Train Size: ', train_data.data.size)

Train Size: 51399

In [12]: 1 # train_data.data

In [13]: 1 train_loader = torch.utils.data.DataLoader(dataset = train_data

```

```

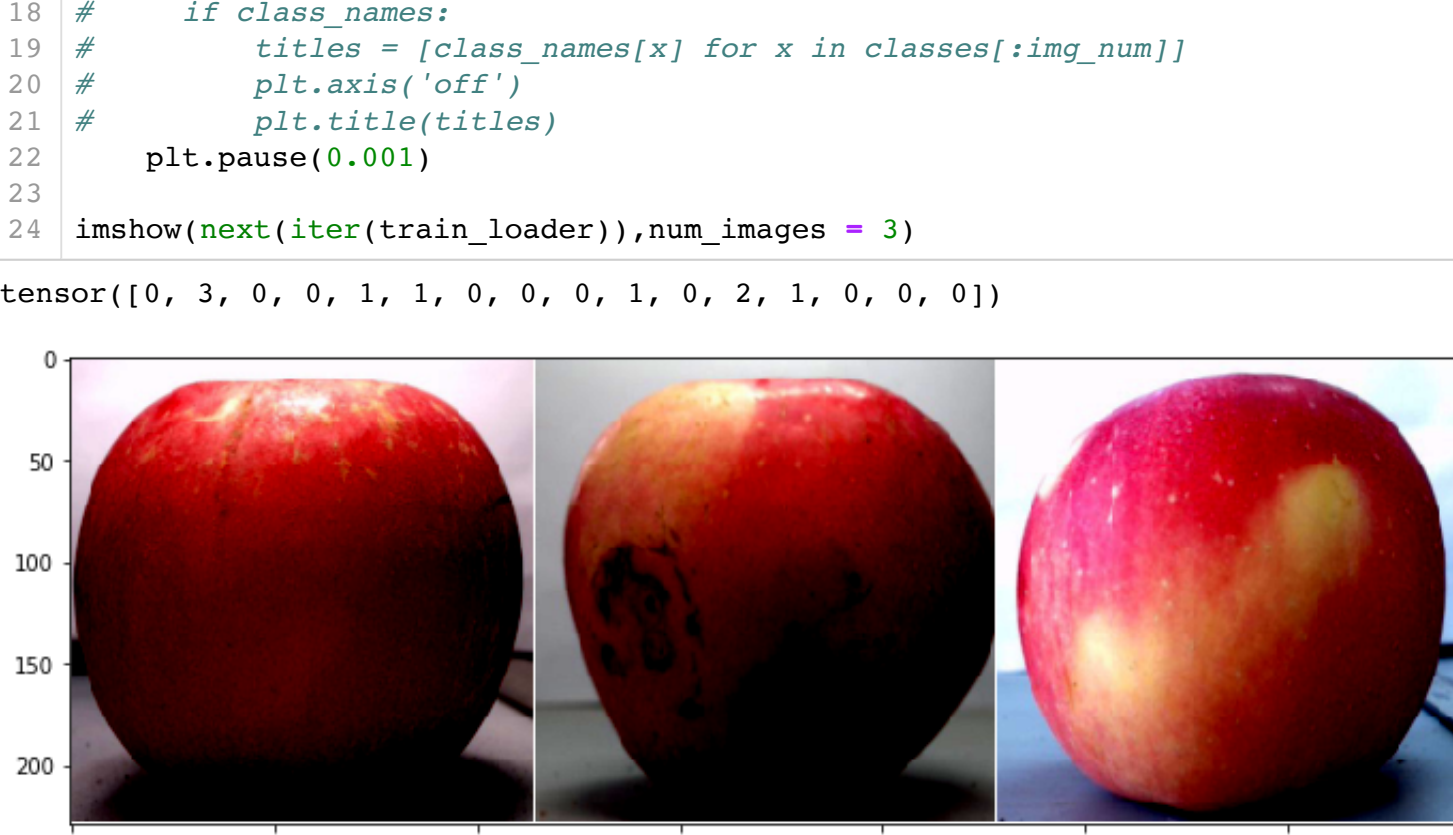
2         batch_size = BATCH_SIZE,
3         shuffle = True)
4     test_loader = torch.utils.data.DataLoader(dataset = test_data,
5         batch_size = BATCH_SIZE)
6
In [14]: 1 def imshow(batch, class_names=None, num_images=4):
2         plt.figure(figsize=(4 * num_images, 4))
3         img, classes, ids, labels = batch
4         print(classes)
5         img_num = min(num_images, img.shape[0])

```

```

6
7     grid = torchvision.utils.make_grid(imgs[:img_num],
8                                         nrow=img_num,
9                                         padding = 1,
10                                        pad_value = 1)
11
12     grid = grid.cpu().numpy().transpose((1, 2, 0))
13     # mean = np.array([0.485, 0.456, 0.406])
14     # std = np.array([0.229, 0.224, 0.225])
15     # grid = std * grid + mean
16     grid = np.clip(grid, 0, 1)
17
18     plt.imshow(grid)

```



```
In [15]: 1 torch.manual_seed(117)

Out[15]: <torch._C.Generator at 0x2af8c5a19a70>
```

Create Model

```
In [16]: 1 class Net(nn.Module):
2         2 def __init__(self):
3             3 self.conv1 = nn.Conv2d(1, 16, kernel_size=5, padding=2)
4             4 self.conv2 = nn.Conv2d(16, 16, kernel_size=5, padding=2)
5             5 self.pool = nn.MaxPool2d(2, 2)
6             6 self.fc1 = nn.Linear(16 * 5 * 5, 120)
7             7 self.fc2 = nn.Linear(120, 10)
8             8 self.softmax = nn.Softmax(dim=-1)
9         9 def forward(self, x):
10            10 x = self.pool(self.conv1(self.conv2(x)))
11            11 x = self.fc1(x.view(-1, 16 * 5 * 5))
12            12 x = self.fc2(x)
13            13 return self.softmax(x)
```

5	
6	

```

8         ('Conv2d_2', nn.Conv2d(64, 192, 5, 1, 2)),
9         ('ReLU_1', nn.ReLU(inplace=True)),
10        ('MaxPool2d_2', nn.MaxPool2d(3, 2, 0)),
11        ('Conv2d_3', nn.Conv2d(192, 384, 5, 1, 1)),
12        ('ReLU_3', nn.ReLU(inplace=True)),
13        ('Conv2d_4', nn.Conv2d(384, 256, 3, 1, 1)),
14        ('ReLU_4', nn.ReLU(inplace=True)),
15        ('Conv2d_5', nn.Conv2d(256, 256, 3, 1, 1)),
16        ('ReLU_5', nn.ReLU(inplace=True)),
17        ('MaxPool2d_3', nn.MaxPool2d(3, 2, 0))
18    ))
19
20    self.dense = torch.nn.Sequential(OrderedDict([
21        ('Dropout_1', nn.Dropout()),
22        ('Linear_1', nn.Linear(9216, 4096)),
23        ('ReLU_1', nn.ReLU(inplace=True)),
24        ('Dropout_2', nn.Dropout()),
25        ('Linear_2', nn.Linear(4096, 4096)),
26        ('ReLU_2', nn.ReLU(inplace=True)),
27        ('Linear_3', nn.Linear(4096, num_classes))
28    ]))
29
30    if init_weights:
31        self._initialize_weights()
32
33    def forward(self, x):
34        x = self.conv(x)
35        x = x.view(x.size(0), -1)
36        x = self.dense(x)
37        return x
38
39    def _initialize_weights(self):
40        for m in self.modules():
41            if isinstance(m, nn.Conv2d):
42                n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
43                m.weight.data.normal_(0, math.sqrt(2. / n))
44                if m.bias is not None:
45                    m.bias.data.zero_()
46            elif isinstance(m, nn.Linear):
47                m.weight.data.normal_(0, 0.01)
48                m.bias.data.zero_()
49

```

```
4 print("Try to use pre-trained model")
5 filePath = os.path.join(DataDir, 'AppleModel.p
```

```

7         print("Loading...")
8         model = torch.load(filePath)
9         print("Success!")
10    else:
11        print("Failed! Create a new model instead.")
12    return model

```

```
3 print("Start train
4 print('epoch: {},
5 model.train()
```

```

6 correct = 0
7 total_loss = 0
8
9 total_acc = 0
10 train_list = {'loss':[], 'acc':[], 'idx':[]}
11
12
13 for batch_idx, (inputs, targets, ids, labels) in enumerate(tqdm_notebook(train_loader, total=len(train_loader),
14                               idx_target = targets.clone(), # clone labels
15                               inputs, targets = Variable(inputs), Variable(targets)
16                               if use_cuda:
17                                   inputs, targets = inputs.to(device), targets.to(device)
18
19     optimizer.zero_grad() # Clears the gradients of all optimized
20     outputs = model(inputs) # output will be batch_size * class_size
21
22     loss = F.cross_entropy(outputs, targets) # calculate loss
23     optimizer.backward()
24     optimizer.step()
25
26     total_loss += loss.data.item() # add loss to total_loss
27
28     pred = outputs.data.max(1)[1] # get the index of the max log-probability
29     correct = pred.cpu().eq(idx_target).sum()
30     total_acc += correct # the amount of correct items
31
32 if batch_idx % 5 == 0 and batch_idx >= 0:
33     acc = float(correct) * 1.0 / len(inputs) * 100
34     train_list['loss'].append(loss.data.item())
35     train_list['acc'].append(acc)
36     train_list['idx'].append(batch_idx)
37     print('Train Epoch: {} [{:5}/{:5}] Loss: {:.6f} Acc: {:.4f} lr: {:.2e}'.format(
38         epoch,
39         batch_idx + len(inputs),
40         len(train_loader.dataset),
41         loss.data.item(),
42         acc,
43         optimizer.param_groups[0]['lr']))
44 print("Train Info: Loss: {:.6f}, Acc: {:.6f}, lr: {:.2e}".format(total_loss,
45                               100 * float(total_acc)/len(train_loader),
46                               optimizer.param_groups[0]['lr']))
47
48 # loss pic
49 plt.plot(train_list['idx'],
50          train_list['loss'],
51          "b", linewidth=1,
52          label = "epoch_" + str(epoch)) #在当前绘图对象绘图 (x轴, y轴, 蓝色虚线, 线宽度)
53 plt.xlabel("iterations") #x轴标签
54 plt.ylabel("loss") #y轴标签
55 plt.title("Lost Analysis") #图标题
56 plt.show()
57
58 #acc pic
59 plt.plot(train_list['idx'],
60          train_list['acc'],
61          "r", linewidth=1,
62          label = "epoch_" + str(epoch)) #在当前绘图对象绘图 (x轴, y轴, 蓝色虚线, 线宽度)
63 plt.xlabel("iterations") #x轴标签
64 plt.ylabel("acc") #y轴标签
65 plt.title("Accuracy Analysis") #图标题
66 plt.show()

```

Create Test Function

```

In [30]: def test(model, device, test_loader, optimizer, epoch):
1         model.eval()
2
3         test_loss = 0
4         correct = 0
5         global best_acc
6         global lr
7         global patience
8
9
10        for batch_idx, (inputs, target, ids, labels) in enumerate(tqdm_notebook(test_loader, total=len(test_loader),
11                               idx_target = target.clone())

```

```
15         inputs =
16
17         if use_cuda:
```

```

18         inputs, target = inputs.to(device), target.to(device)
19
20         output = model(inputs)
21         test_loss += F.cross_entropy(output, target).data.item()
22         pred = output.data.max(1)[1] # get the index of the max log-probability
23         correct += pred.eq(target.view_as(pred)).sum().item()
24         test_loss_percent = test_loss / len(test_loader) # average over number of mini-batch
25         acc = 100. * float(correct) / len(test_loader.dataset)
26
27
28     print('Test set: Average loss: {:.4f}, Accuracy: {}/{} ({:.2f}%)\n'.format(

```

```

30         correct,
31         len(test_loader.dataset),
32         acc))
33
34     if acc > best_acc:
35         save_file = True # config. If need to save model
36
37         patience = patience + 1 # if the acc of the model better than older one two times, than update it.
38         print("Best acc is {},lower than current acc {}. Patience set to {}".format(best_acc,acc,patience))
39
40     if patience == 2:

```

43
44
45

```

46         lr_scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'max', patience=10)
47         new_file = os.path.join(DataDir, 'AppleModel.pth')
48         torch.save(model, new_file)
49         print("Save a new model to {}".format(new_file))
50     best_acc = acc
51

```

Create model

```

In [30]: 1 model = AppleModel(pretrained = True, num_classes = N_CLASSES, init_weights = True)

```

```
2 if use_cuda:
3     model.to(device)
4     print(model)

Try to use pre-trained model
Loading...
Success!
Net(
  (conv): Sequential(
    (conv2d_1): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (ReLU_1): ReLU(inplace)
    (maxPool2_1): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```
(Conv2d_2: Conv2d[64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2)]
    ReLU(inplace)
    (MaxPool2d_2: MaxPool2d[kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False]
        Conv2d_3: Conv2d[192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)]
        ReLU(inplace)
        Conv2d_4: Conv2d[384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)]
        ReLU(inplace)
        Conv2d_5: Conv2d[256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)]
        ReLU(inplace)
        (MaxPool2d_3: MaxPool2d[kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False])
    )
)
```



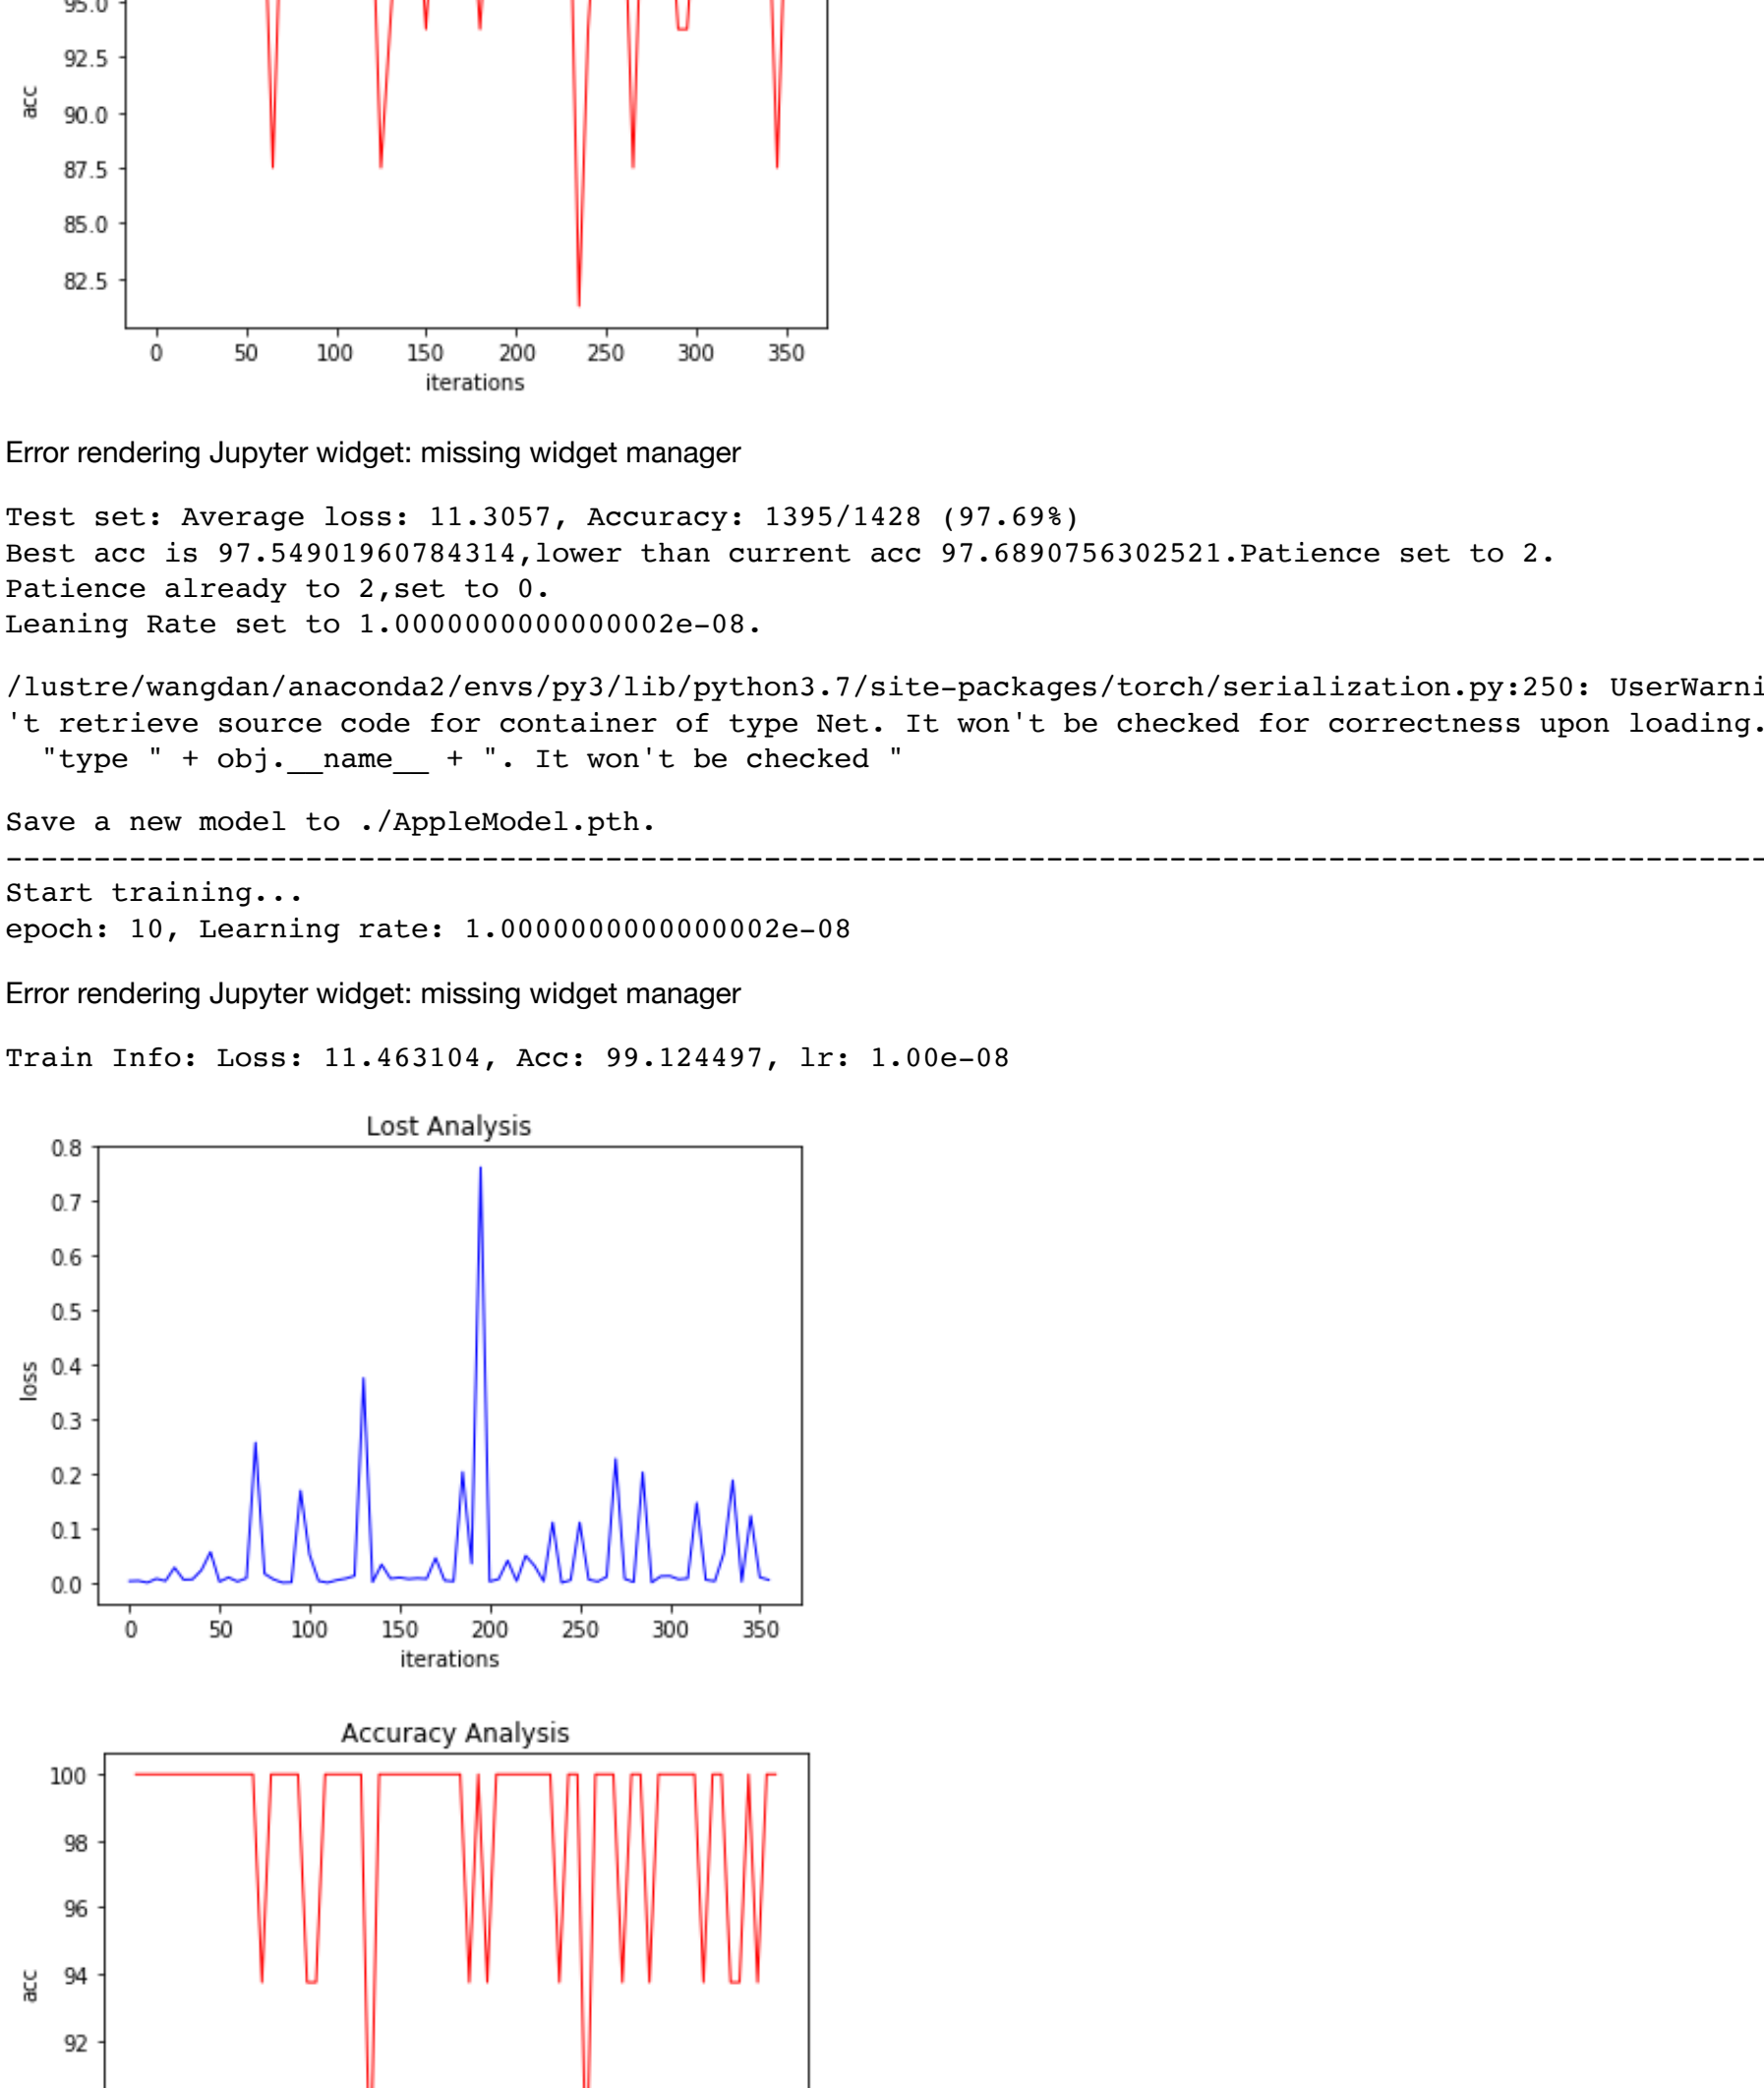

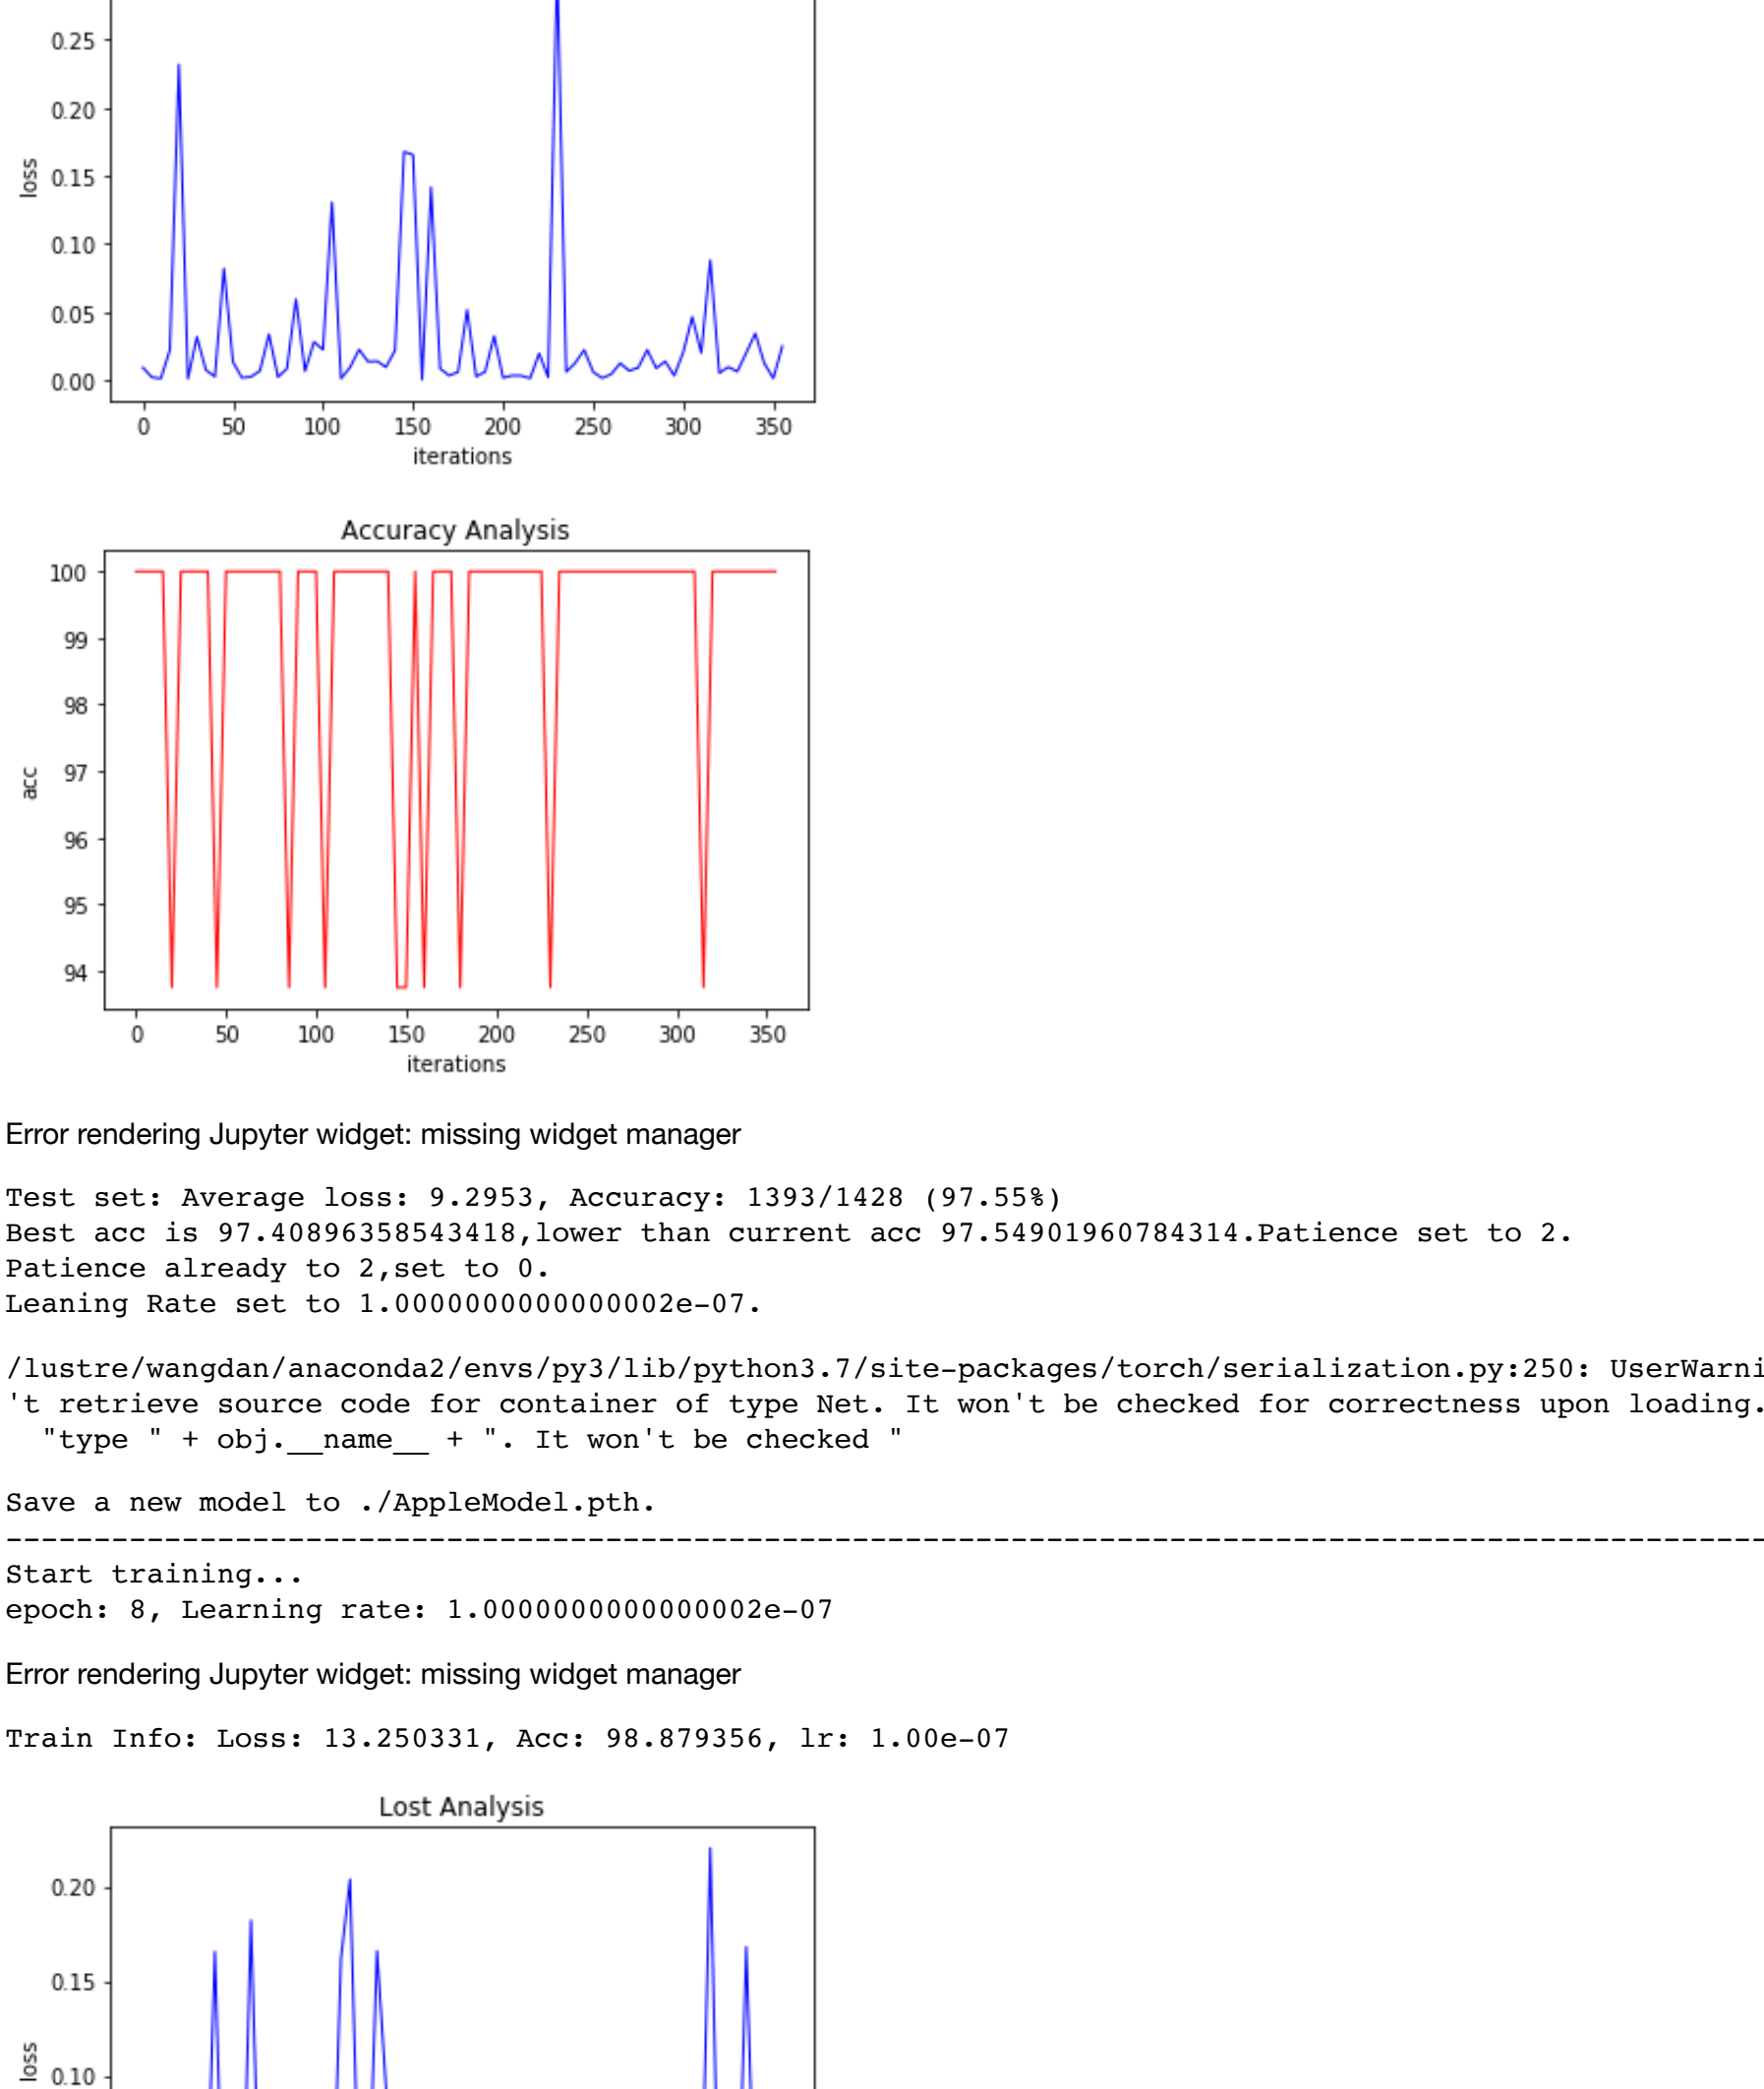

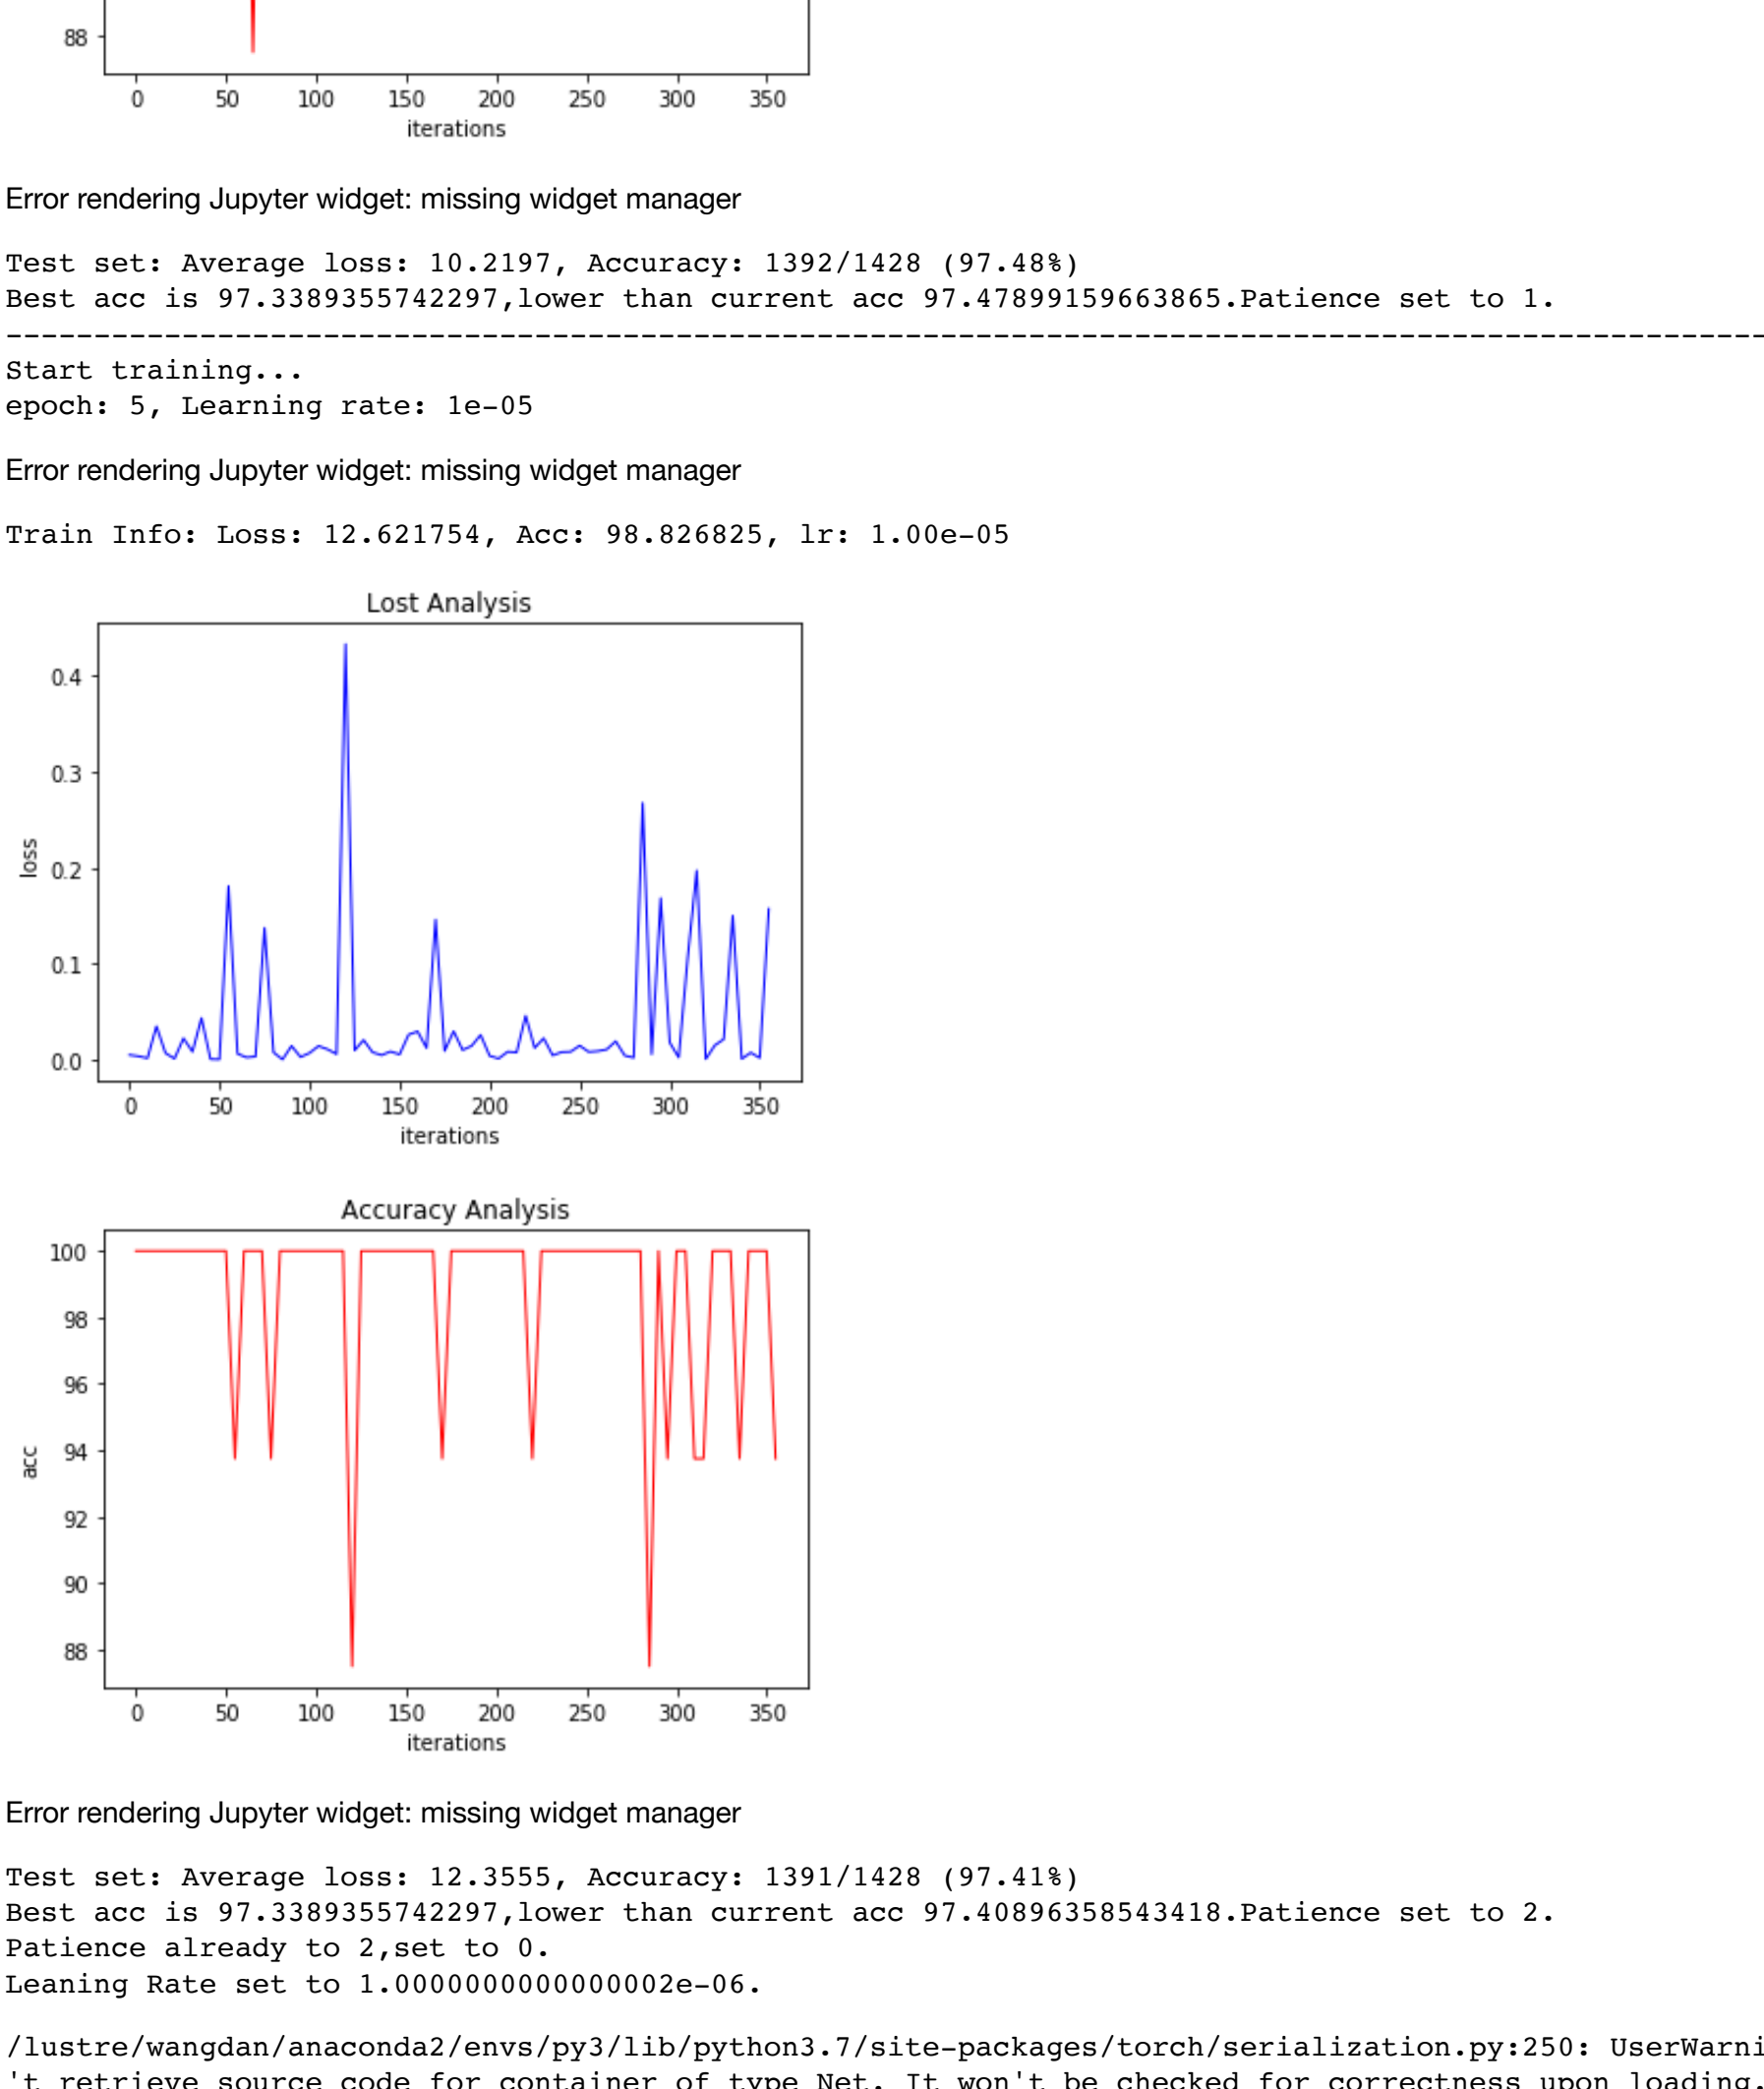
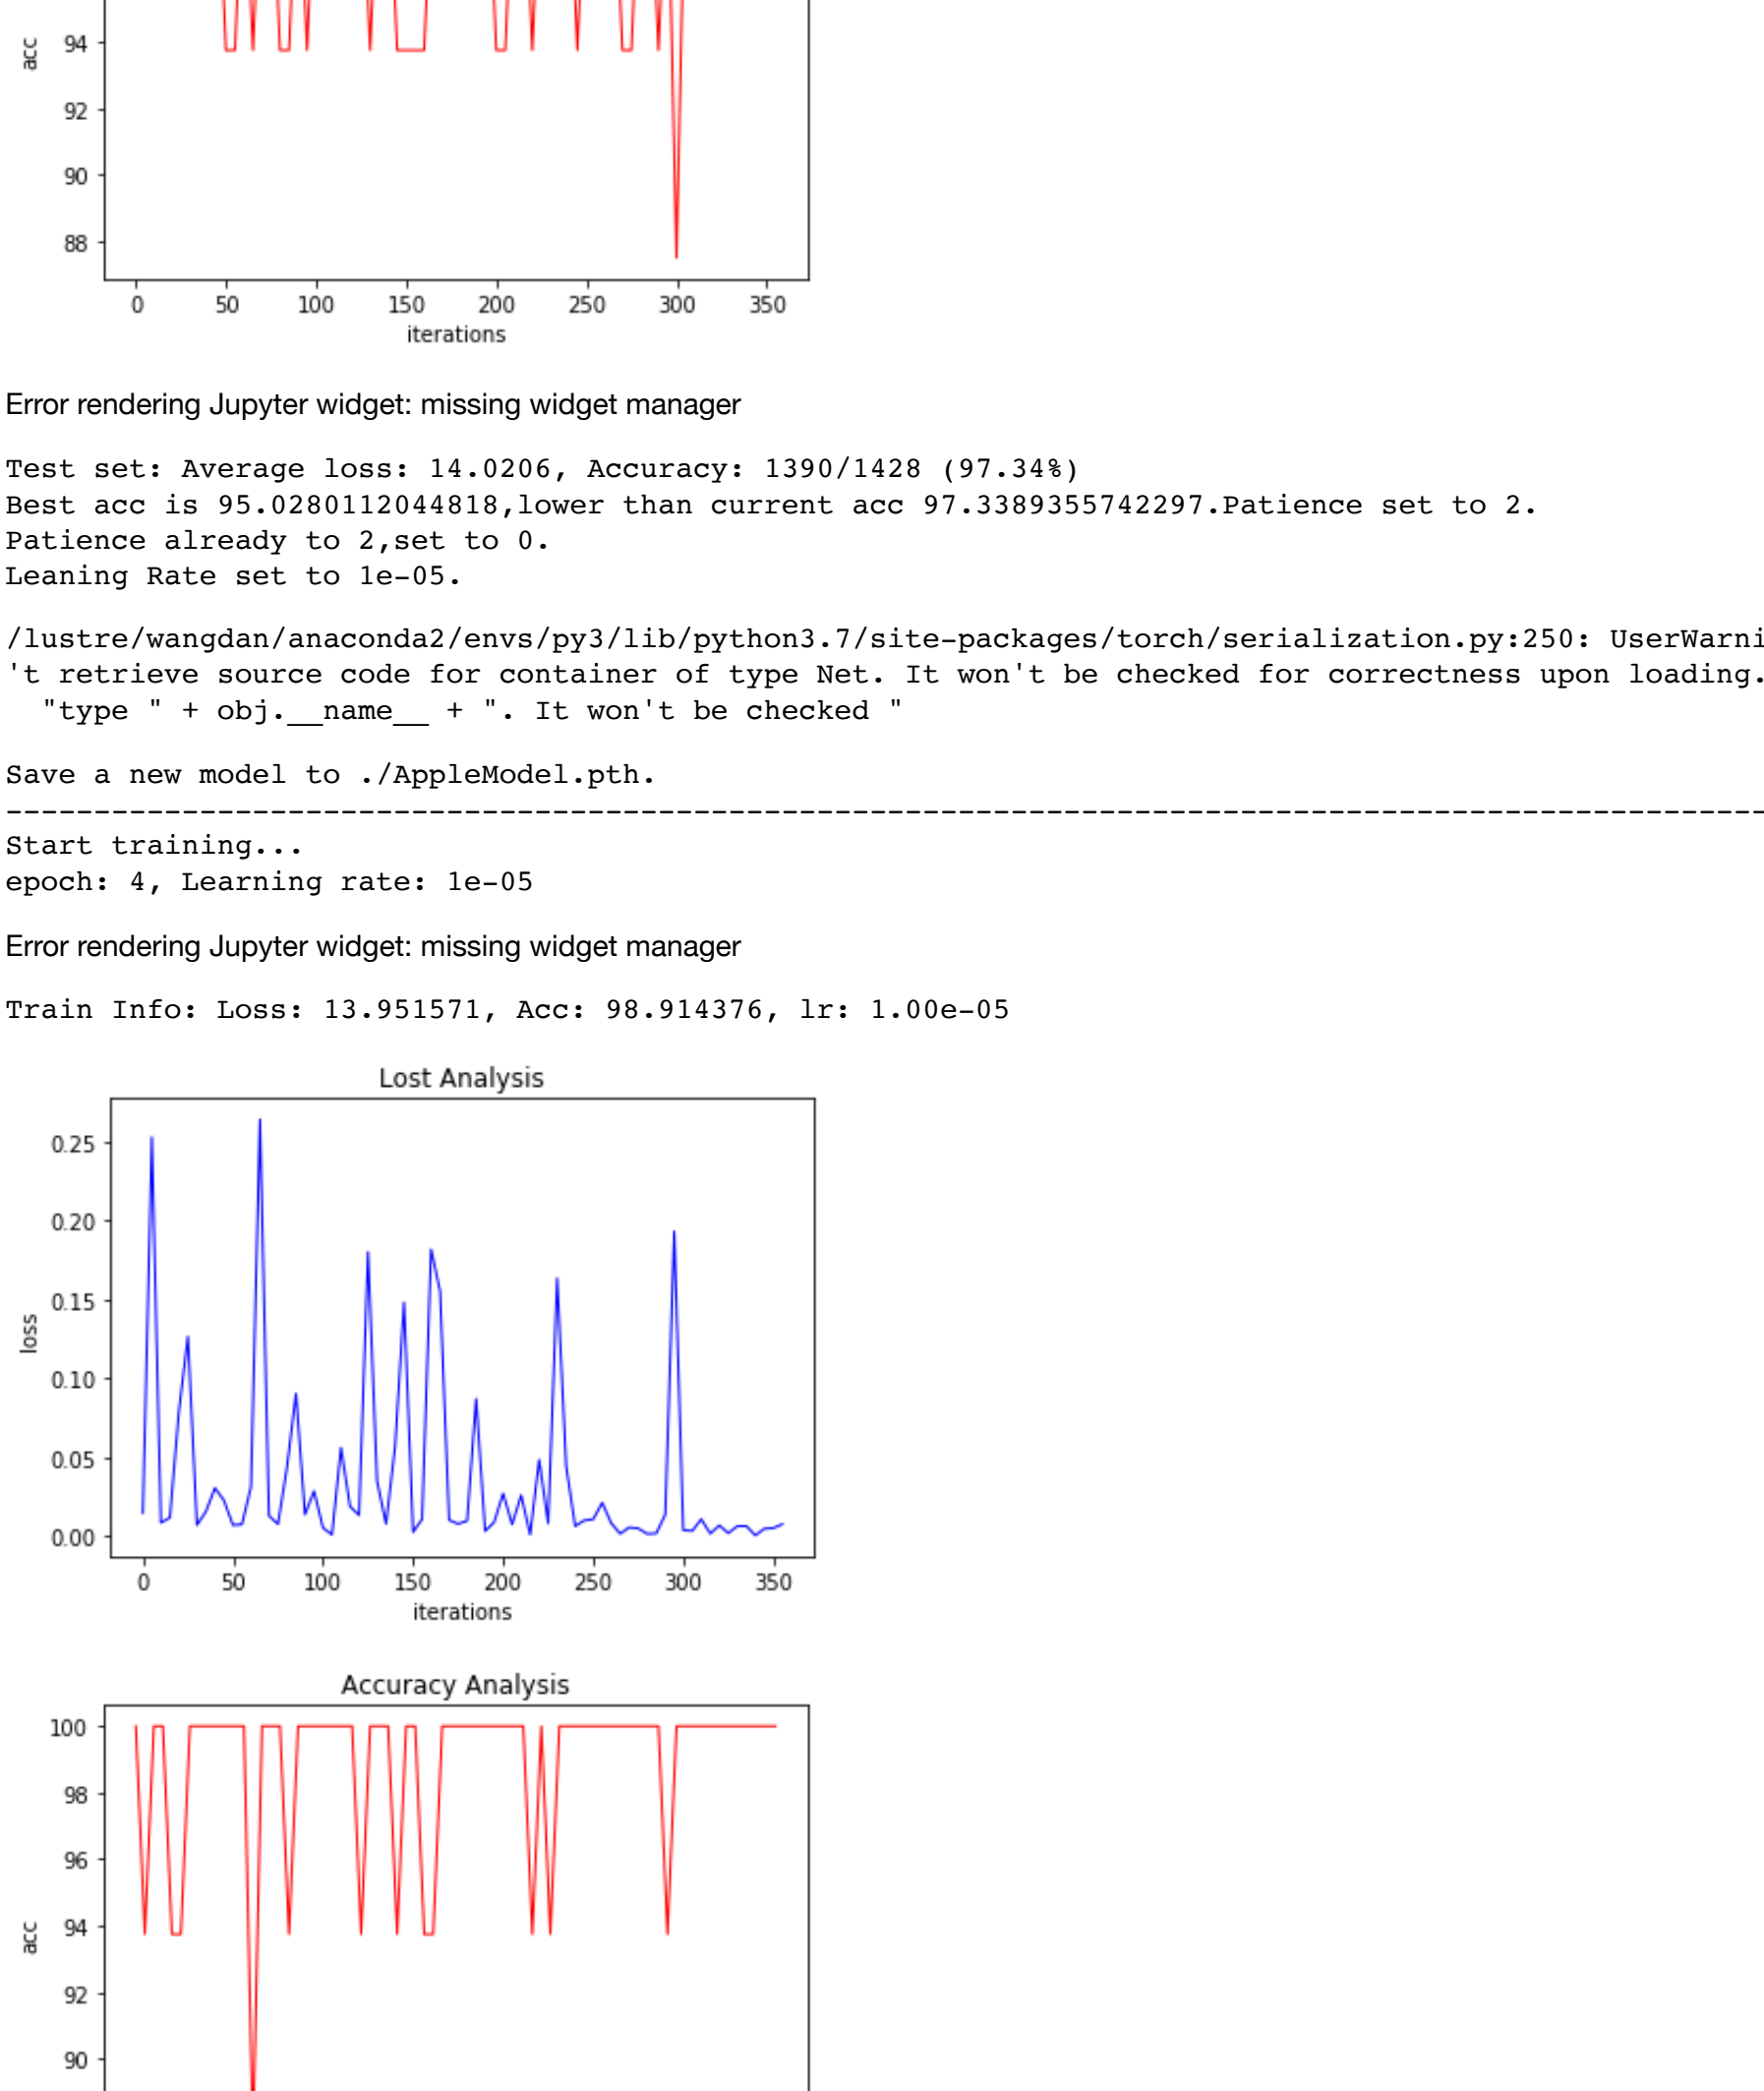
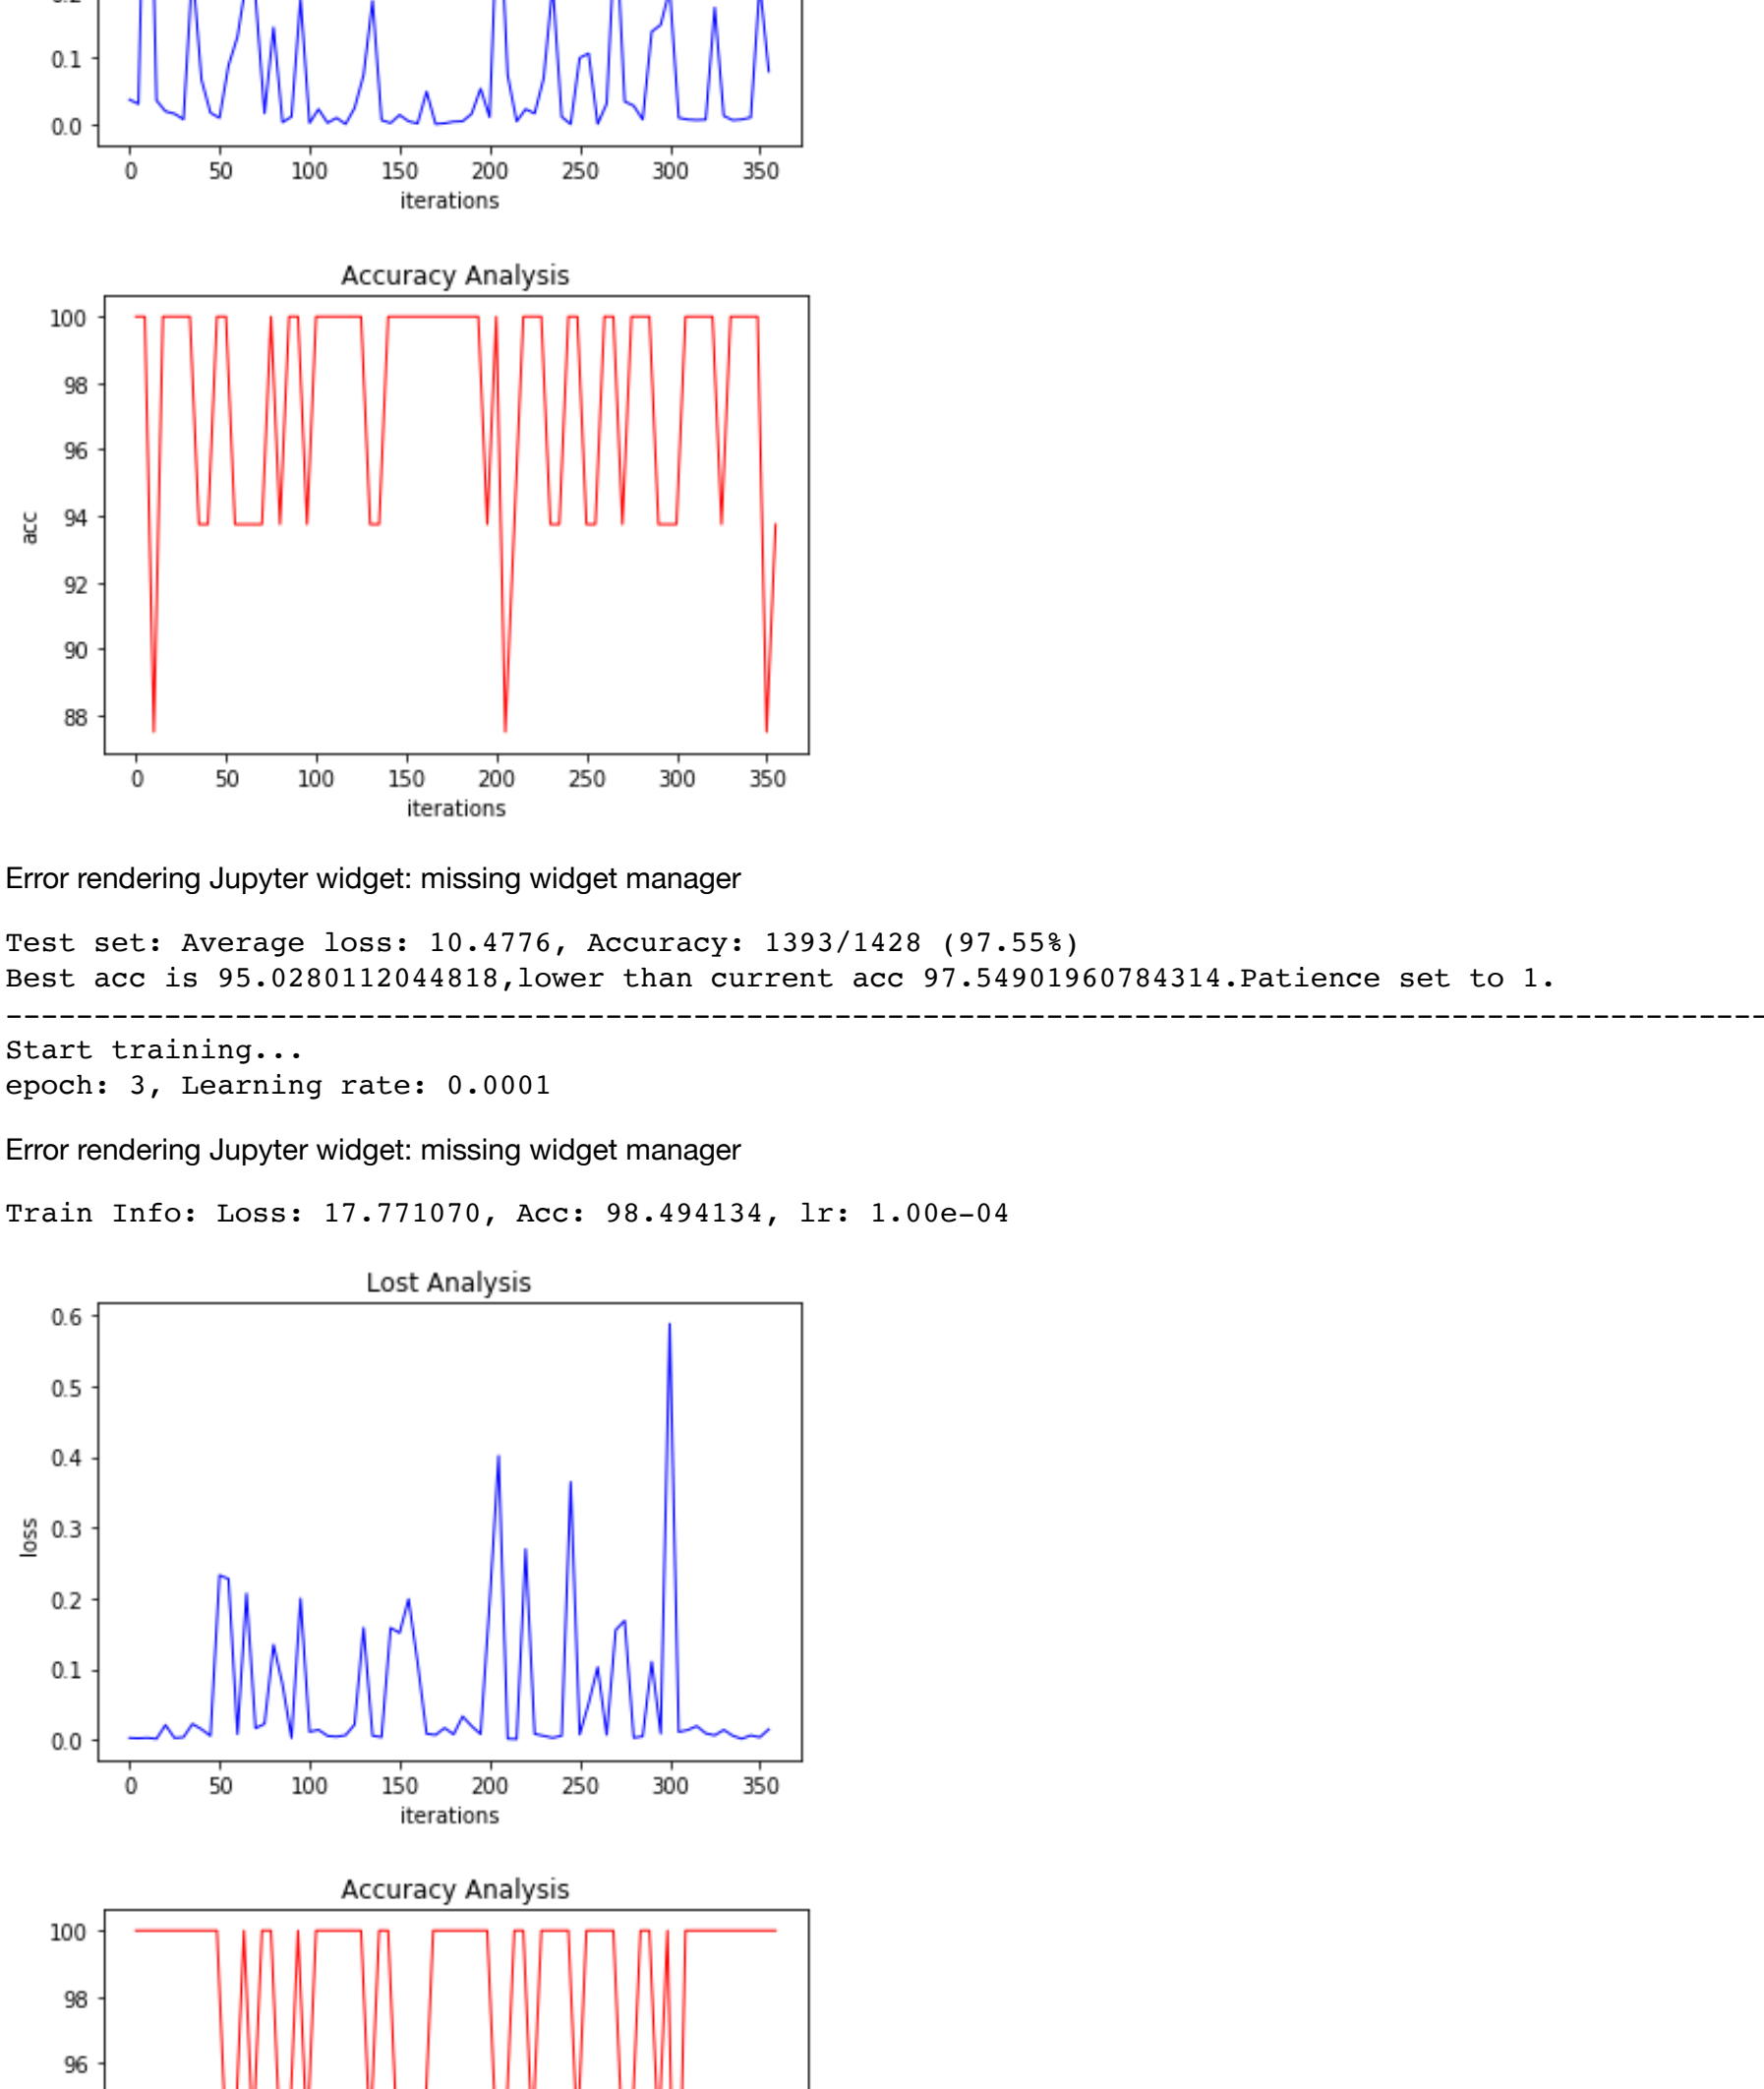
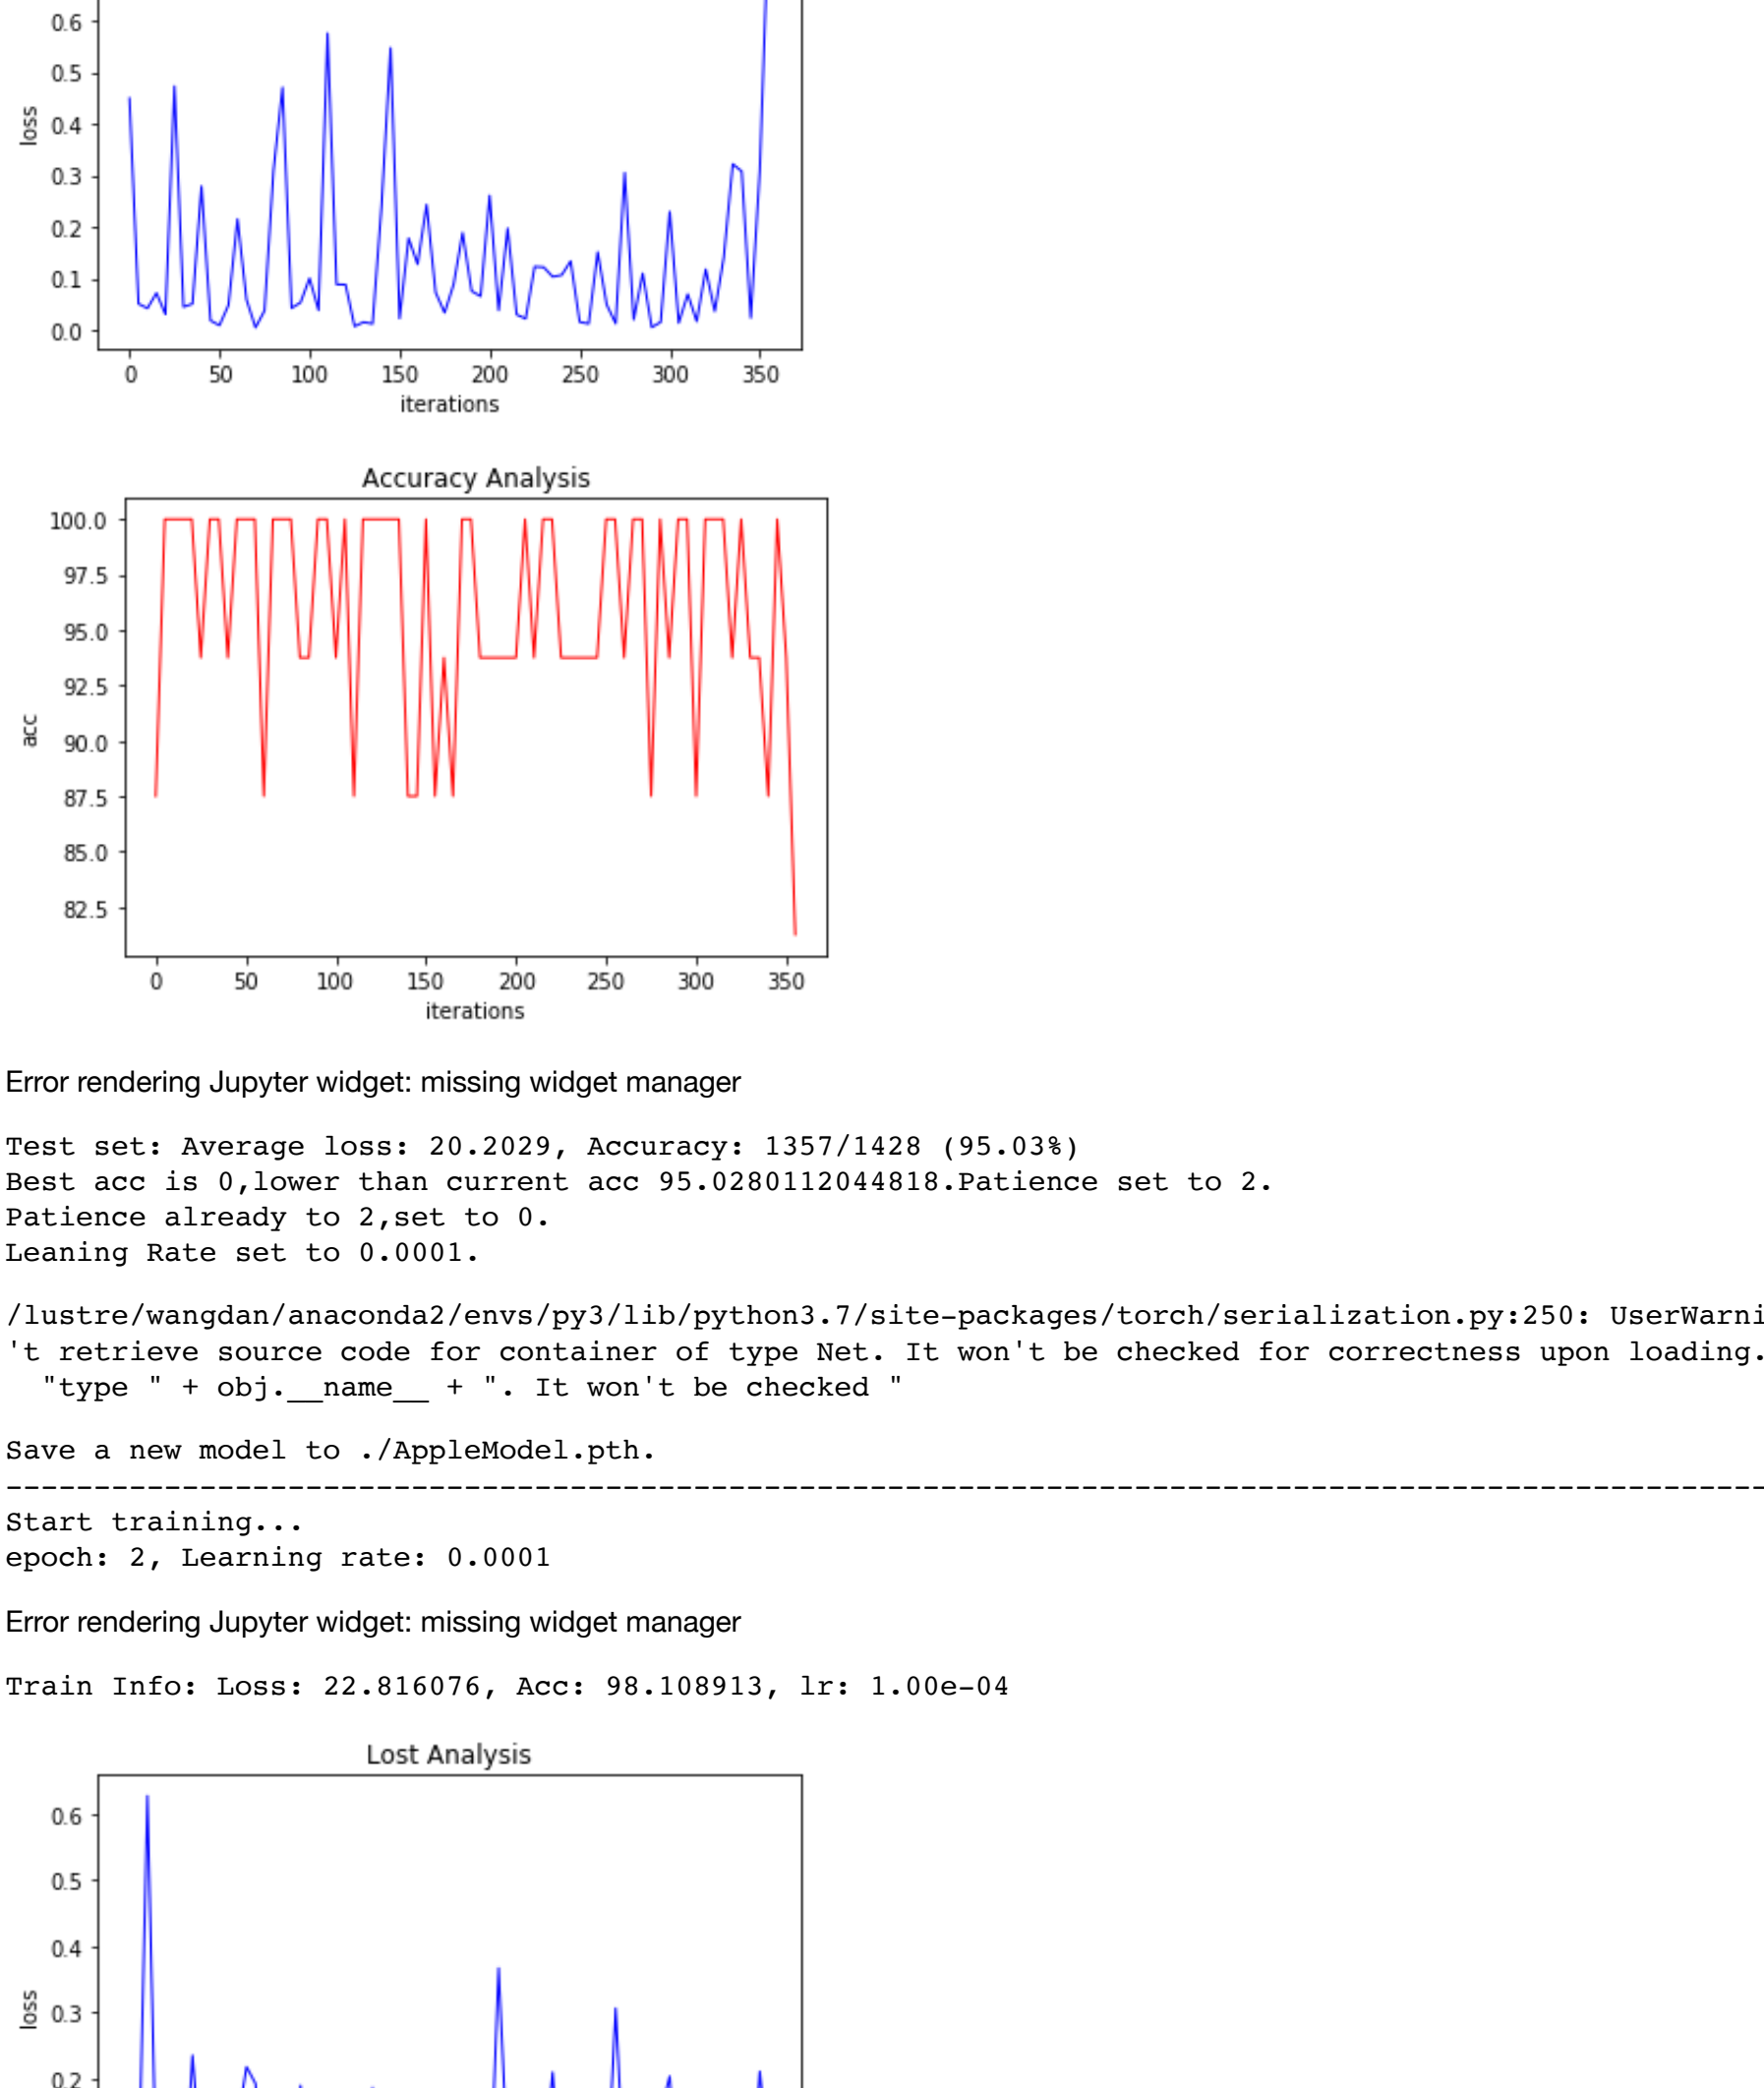

```
(Dropout_1): Dropout(p=0.5)
  (Linear_1): Linear(in_features=9216, out_features=4096, bias=True)
    (ReLU_1): ReLU(inplace)
  (Dropout_2): Dropout(p=0.5)
    (Linear_2): Linear(in_features=4096, out_features=4096, bias=True)
      (ReLU_2): ReLU(inplace)
    (Linear_3): Linear(in_features=4096, out_features=4, bias=True)
  )
)
```



```
In [48]: 1 best_acc = 0
2 lr = 0
3 patience = 0
4
5 try:
6     epoch_size = 30
7     for epoch in range(epoch_size):
8         if epoch == 0:
9             lr = 0.001 # init learning rate
```

```

11         optimizer = optim.Adam(model.parameters(), lr = lr, weight_decay=5e-4)
12
13     train(model, device, train_loader, optimizer, epoch)
14     test(model, device, test_loader, optimizer, epoch)
15
16     torch.cuda.empty_cache()
17 except Exception as e:
18     import traceback
19     traceback.print_exc()
20 finally:
21     print("Best Result: {:.3f}".format(best_acc))

```



Error rendering Jupyter widget: missing widget manager

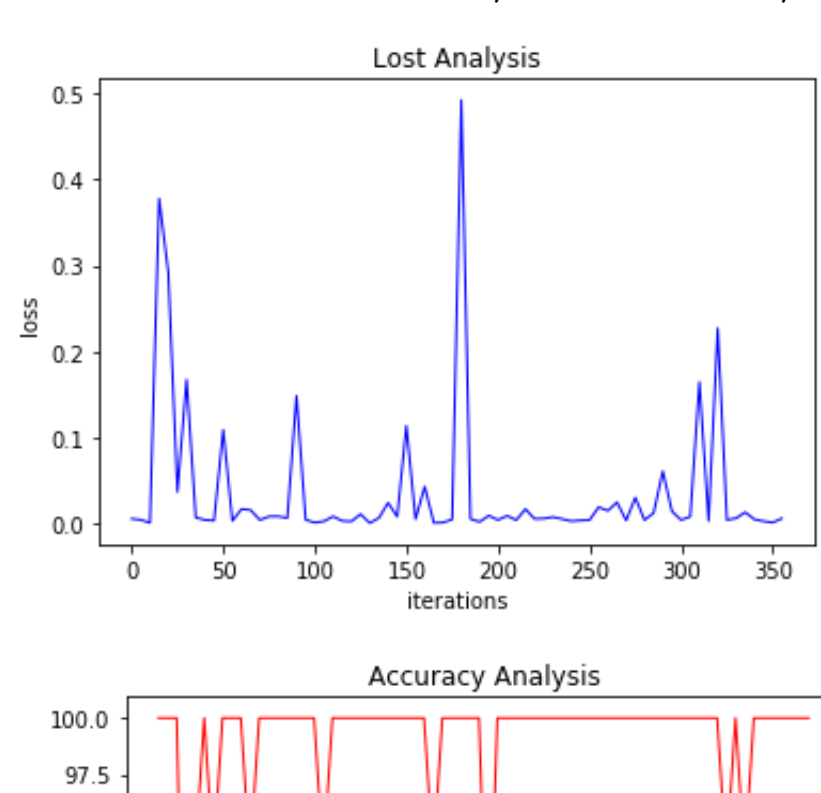
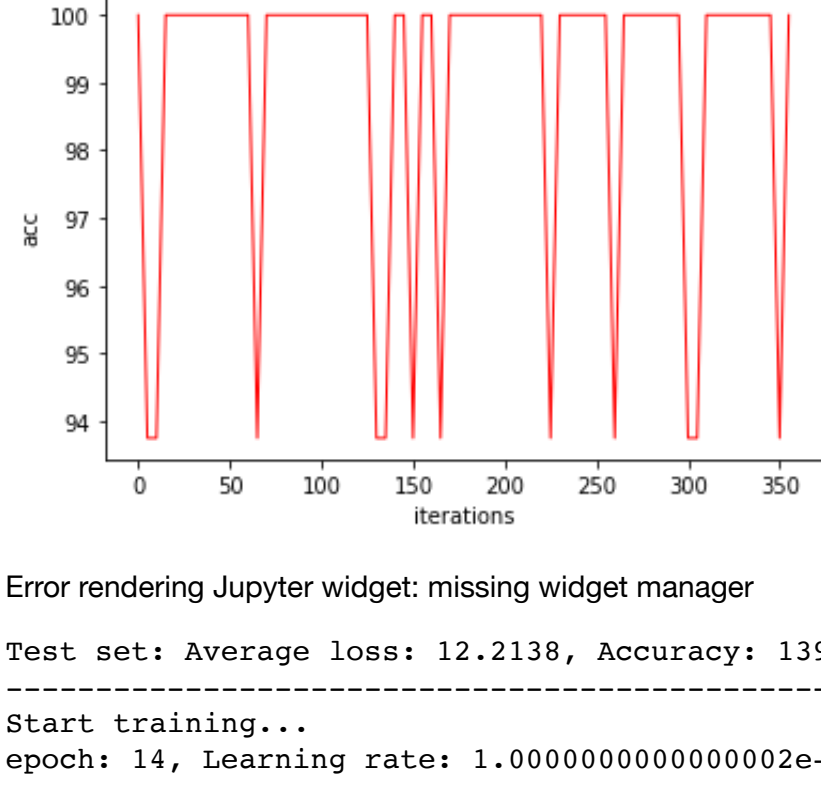
Test set: Average loss: 11.2170, Accuracy: 1394/1428 (97.62%)

Start training...

epoch: 13, Learning rate: 1.0000000000000002e-08

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.242456, Acc: 98.896866, lr: 1.00e-08



Error rendering Jupyter widget: missing widget manager


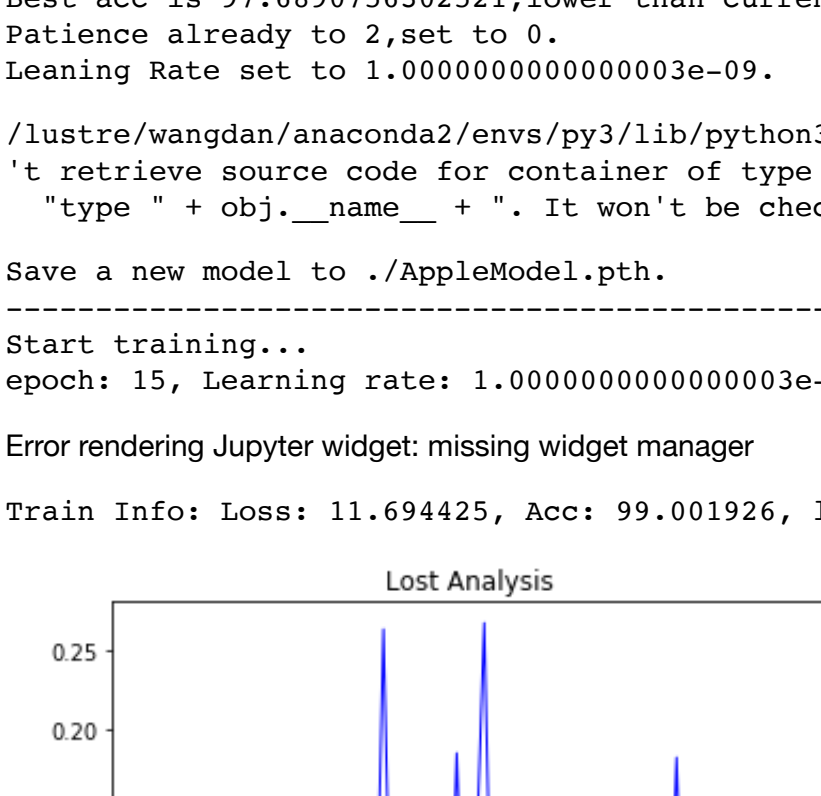
Test set: Average loss: 12.2138, Accuracy: 1391/1428 (97.41%)

Start training...

epoch: 14, Learning rate: 1.0000000000000002e-08

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.301325, Acc: 98.931886, lr: 1.00e-08



Error rendering Jupyter widget: missing widget manager

Test set: Average loss: 10.5185, Accuracy: 1397/1428 (97.83%)
Best acc is 97.6890756302521,lower than current acc 97.82913165266106.Patience set to 2.

Patience already to 2,set to 0.

Learning Rate set to 1.0000000000000003e-09.

/lustre/wangdan/anaconda2/envs/py3/lib/python3.7/site-packages/torch/serialization.py:250: UserWarning: Could not retrieve source code for container of type Net. It won't be checked for correctness upon loading.
"type " + obj.__name__ + ". It won't be checked"



Save a new model to ./AppleModel.pth.

Start training...

epoch: 15, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.694425, Acc: 99.001926, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager


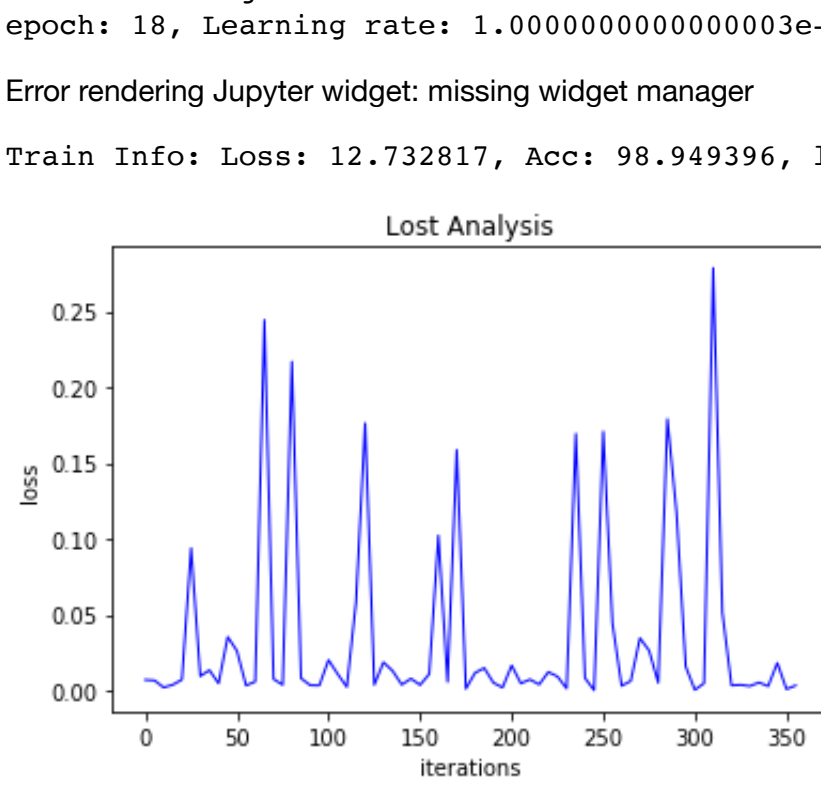
Test set: Average loss: 10.4531, Accuracy: 1396/1428 (97.76%)

Start training...

epoch: 16, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.083903, Acc: 98.914376, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager

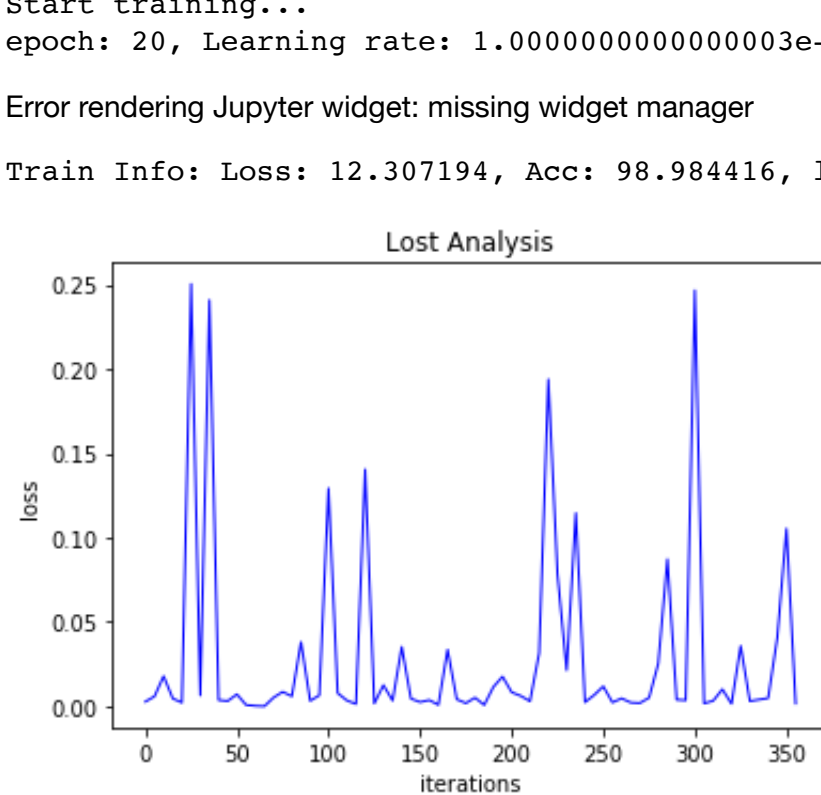
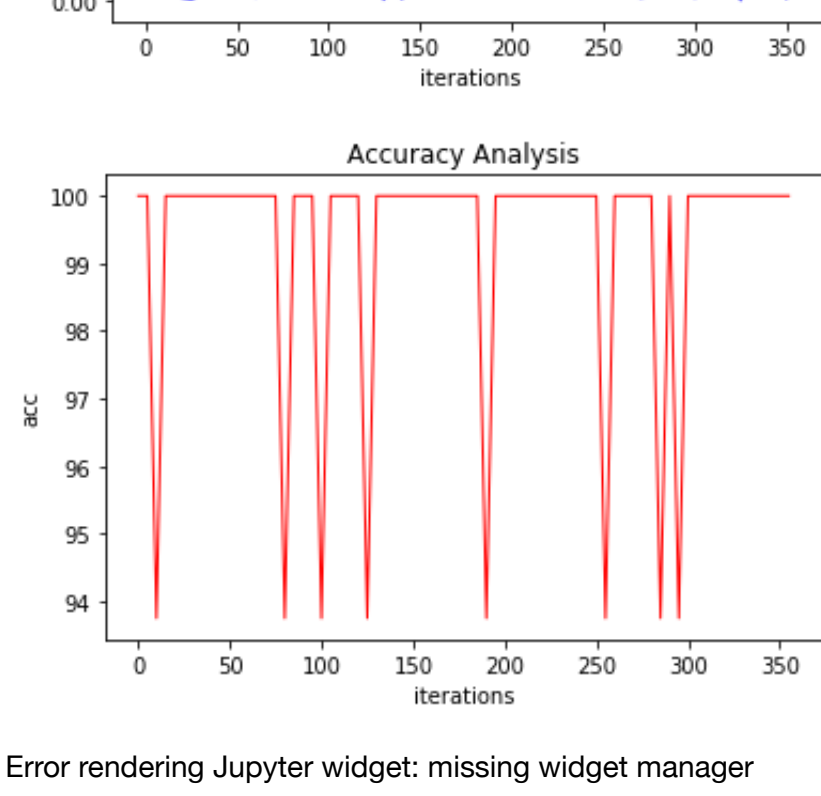
Test set: Average loss: 12.0778, Accuracy: 1391/1428 (97.41%)

Start training...

epoch: 17, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.707259, Acc: 98.914376, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager

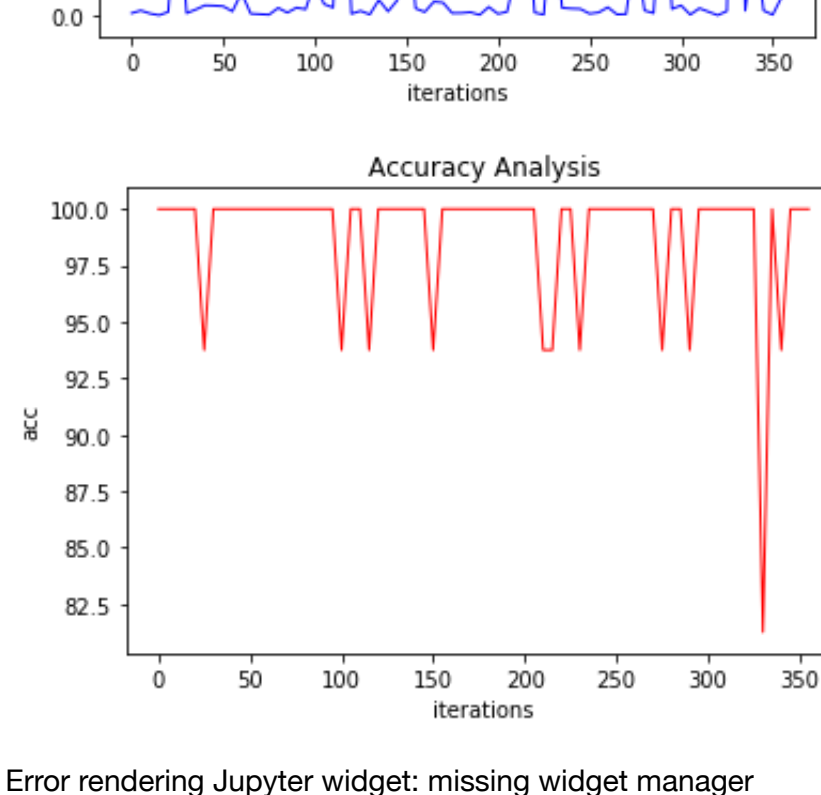
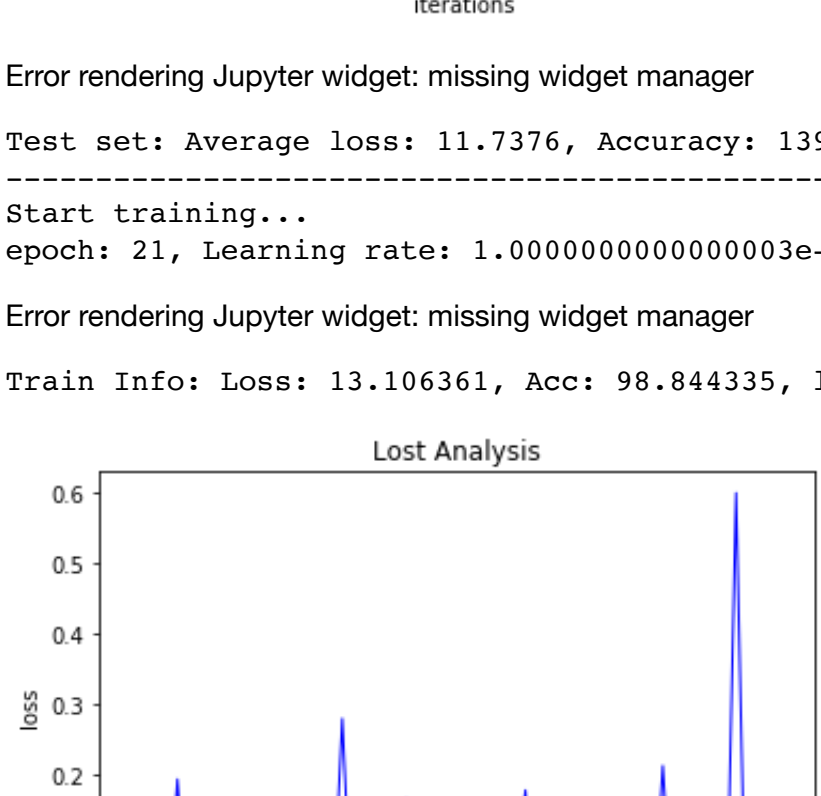
Test set: Average loss: 10.2145, Accuracy: 1396/1428 (97.76%)

Start training...

epoch: 18, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.732817, Acc: 98.949396, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager

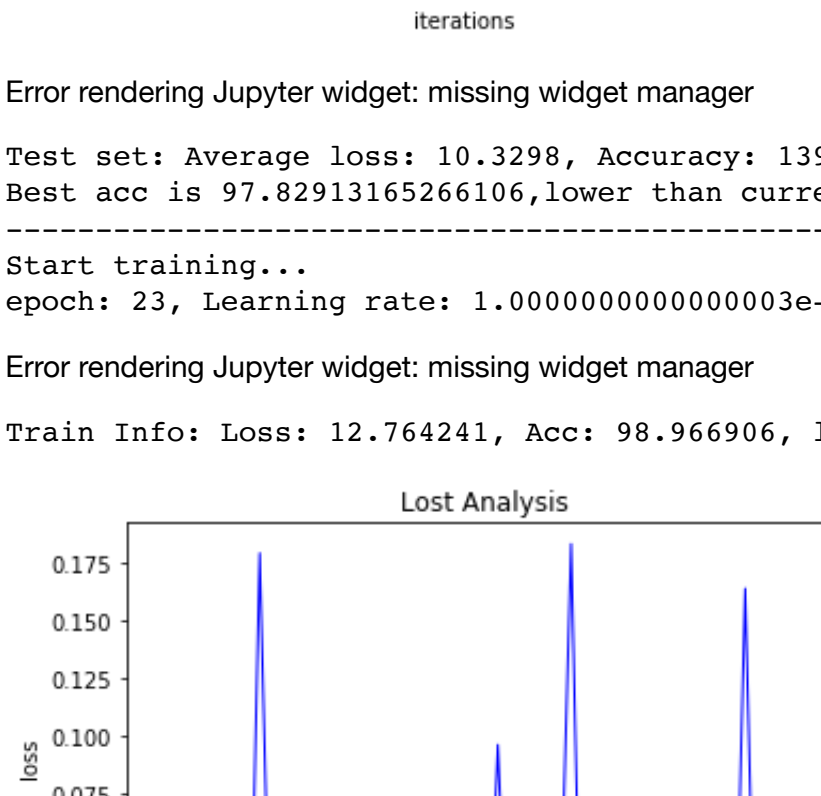

Test set: Average loss: 11.6883, Accuracy: 1395/1428 (97.69%)

Start training...

epoch: 19, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.701669, Acc: 99.019436, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager


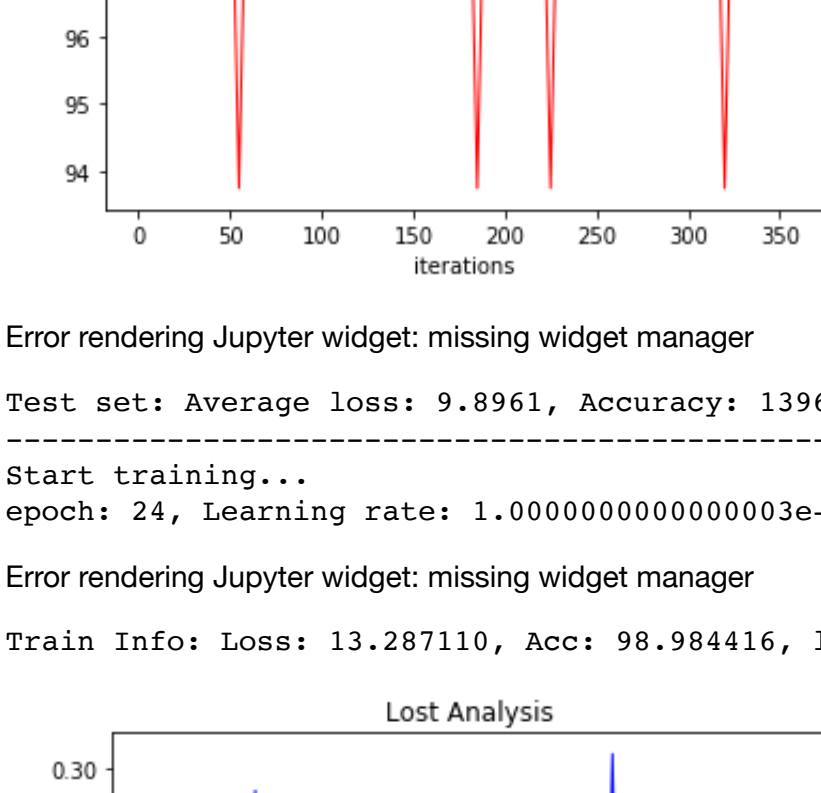
Test set: Average loss: 12.1495, Accuracy: 1393/1428 (97.55%)

Start training...

epoch: 20, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.307194, Acc: 98.844316, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager


Test set: Average loss: 11.7376, Accuracy: 1395/1428 (97.69%)

Start training...

epoch: 21, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 13.106361, Acc: 98.844335, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager

Test set: Average loss: 11.2668, Accuracy: 1394/1428 (97.62%)

Start training...

epoch: 22, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.581755, Acc: 99.036946, lr: 1.00e-09

Error rendering Jupyter widget: missing widget manager

Test set: Average loss: 10.3298, Accuracy: 1398/1428 (97.90%)

Best acc is 97.82913165266106,lower than current acc 97.89915966386555.Patience set to 1.

Start training...

epoch: 23, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.764241, Acc: 98.966906, lr: 1.00e-09

Error rendering Jupyter widget: missing widget manager

Test set: Average loss: 9.8961, Accuracy: 1396/1428 (97.76%)

Start training...

epoch: 24, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.287110, Acc: 98.984614, lr: 1.00e-09

Error rendering Jupyter widget: missing widget manager

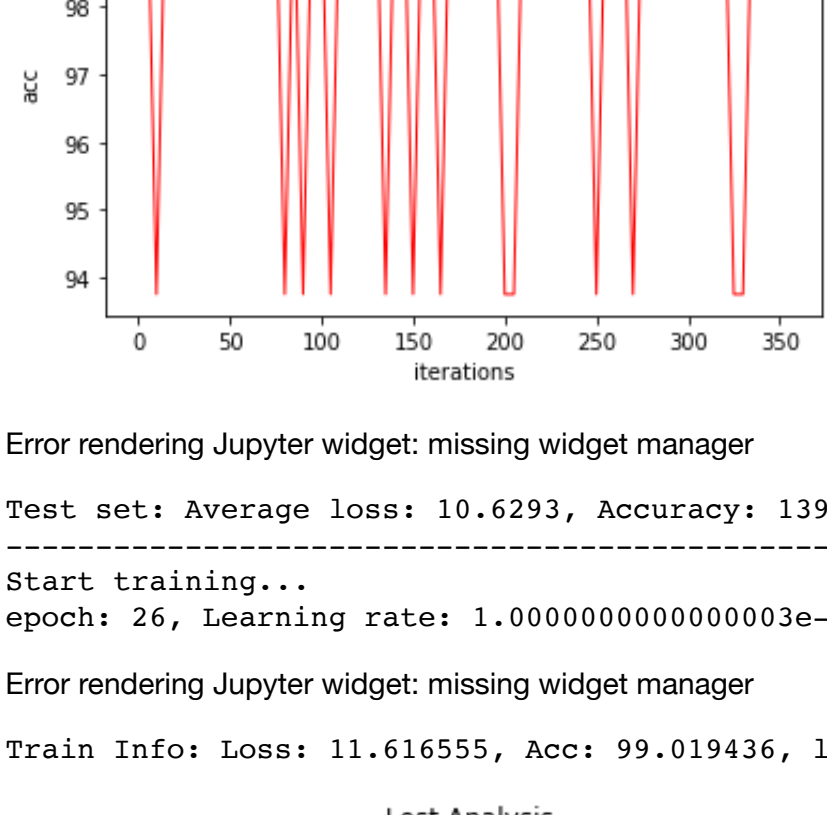

Test set: Average loss: 11.7371, Accuracy: 1395/1428 (97.69%)

Start training...

epoch: 25, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.550140, Acc: 98.966906, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager



Test set: Average loss: 10.6293, Accuracy: 1396/1428 (97.76%)

Start training...

epoch: 26, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.616555, Acc: 99.019436, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager



Test set: Average loss: 10.8890, Accuracy: 1394/1428 (97.62%)

Start training...

epoch: 27, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.734734, Acc: 98.914376, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager


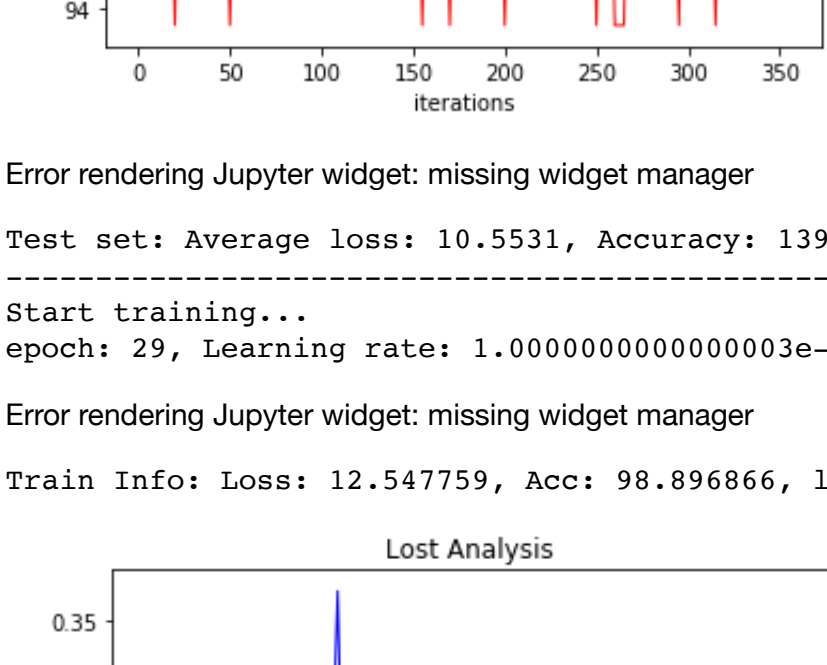
Test set: Average loss: 11.5396, Accuracy: 1392/1428 (97.48%)

Start training...

epoch: 28, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 11.844354, Acc: 98.879356, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager


Test set: Average loss: 10.5531, Accuracy: 1396/1428 (97.76%)

Start training...

epoch: 29, Learning rate: 1.0000000000000003e-09

Error rendering Jupyter widget: missing widget manager

Train Info: Loss: 12.547759, Acc: 98.896866, lr: 1.00e-09



Error rendering Jupyter widget: missing widget manager

Test set: Average loss: 12.7612, Accuracy: 1394/1428 (97.62%)

Best Result: 97.829%

In []: 1

In []: 1